

Cache

Motivation

- ▶ Let's consider a typical “high end” PC...
- ▶ I went to Alienware's site¹ and pulled up their highest priced gaming system
- ▶ Here's its specs:
 - ▶ AMD Ryzen™ Threadripper 2950X (16-Core, 40MB Cache, 4.4GHz Boost)
 - ▶ 64GB Quad Channel HyperX™ DDR4 XMP at 2933MHz
- ▶ Only \$7999! Such a bargain!

¹<https://www.dell.com/en-us/shop/dell-desktop-computers/alienware-area-51-threadripper-edition-gaming-desktop/spd/alienware-area51-r7/dpcwCfxR7maxh>

Explanation

- ▶ So what do those features mean?
 - ▶ The CPU is 4.4GHz
 - ▶ That means one clock tick every 227ps
 - ▶ Picosecond = one trillionth of a second
 - ▶ The memory is DDR4-2933MHz
 - ▶ That means $\sim 23500 \text{ MB/sec} = 23.5 \text{ GB/sec}$

Calculations

- ▶ I found that the average x86-64 instruction is 4 bytes long
- ▶ So: CPU requires $3.6 \times 10^9 \times 4$ bytes per second just to be supplied with enough instructions to be kept constantly busy
 - ▶ That's about 14.4GB/sec
- ▶ But that's not the whole story...

Calculations

- ▶ Many instructions read or write memory
- ▶ I found that on average about 40% of them do
 - ▶ For simplicity, let's assume each access is 32 bits
 - ▶ Many will be 64 bits, so we're being optimistic
- ▶ So we need another $0.4 \times 3.6 \times 10^9 \times 4$ bytes per second
- ▶ Our running total: 20.16GB/sec
- ▶ But that's not the whole story...

Calculations


- ▶ Suppose we have a 16 core CPU
- ▶ Then we need to feed *sixteen* cpu's with $20.16\text{GB/sec} = 322\text{GB/sec}$
 - ▶ Clearly, this outstrips our memory subsystem's speed (23.5GB/sec)
- ▶ What to do?

Channels

- ▶ One trick: When reading, fetch data from two (or four) memory modules simultaneously
 - ▶ Called “channels”: One channel per module
- ▶ When CPU fetches data, read from *both* channels in parallel
- ▶ This doubles (or quadruples) throughput
- ▶ But that's not the whole story...

Latency


- ▶ We face a problem: Latency
- ▶ Imagine a 4GB memory module
- ▶ We *could* structure it internally like so:



- ▶ But that means 4 billion wires from memory to bus. Yuck!

Better

- ▶ Memory chips are internally designed as a grid structure:



- ▶ This means we have $64K \times 64K$ cells
- ▶ That's easier to manage than single row of 4 billion

Better Yet


- ▶ We can arrange memory in several *banks*
 - ▶ Ex: For 4GB of memory: Have 8 banks of 512MB each
 - ▶ Each bank has 16K x 32K cells

Problem

- ▶ Selecting row/column/bank takes time
- ▶ For DDR-2900, about 55ns to initiate transfer
- ▶ That means the CPU ticks almost 200 times before data begins streaming in
- ▶ But that's not the whole story...

Refresh

- ▶ DRAM gets incredible densities
 - ▶ Today, you can get a 64 *billion* byte module that's just 8.2cm × 3cm (about 3"×1¼")
 - ▶ Each DRAM bit is extremely simple: Just a capacitor and a transistor




Problem

- ▶ Capacitors are inherently “leaky”
 - ▶ Charge dissipates over time
 - ▶ DRAM circuitry must periodically “refresh” memory cells
 - ▶ They’re unavailable for access when that’s happening
- ▶ A typical refresh interval might be once every $4\mu\text{s}$
- ▶ A refresh might take around 50ns
- ▶ Remember: CPU ticks every 250ps , so 50ns is a significant amount of time
 - ▶ Memory controller can sometimes “delay” refreshes a bit if CPU wants to access a given region of memory

Solution

- ▶ We'd like to address problems of DRAM memories
- ▶ One possibility: Use SRAM
- ▶ Drawback: More expensive (physically larger)
 - ▶ So we can't get as much capacity in given physical space




Cache

- ▶ Def: (dictionary.com): “A hiding place...for ammunition, food, treasures, etc.”
- ▶ A CPU cache is similar
 - ▶ Small capacity chunk of memory that's fast
- ▶ If we arrange for cache to (nearly) always have the data we need:
We can run at (nearly) full CPU speed


Hierarchy

- ▶ Often, caches are arranged in a hierarchy
 - ▶ Level 1 (L1), Level 2 (L2), Level 3 (L3)
 - ▶ As caches get larger, they usually get slower



Hierarchy

- ▶ Example: Each core has its own L1, L2 cache
- ▶ L3 cache is shared by all cores



Structure

- ▶ Some architectures have $L1 \subseteq L2 \subseteq L3$
- ▶ Others make them disjoint
 - ▶ If data is in L1, it's not in L2
 - ▶ And if data is in L2, it's not in L3

Example

- ▶ Intel Skylake architecture
 - ▶ L1 = 32KB, 4 cycle latency
 - ▶ L2 = 256KB, 12 cycles
 - ▶ L3 = 2MB per core, 44 cycles

Challenge

- ▶ Need to try to keep cache hit rate high
- ▶ Modern CPU's manage cache in the background
 - ▶ So we don't explicitly interact with it
 - ▶ But: Must understand how it works so we can best use it

Design

- ▶ Recall: There's a startup cost for pulling data from memory
- ▶ So: We'd like to spread that cost out as much as possible
- ▶ Thus, fundamental quantum for cache is not bytes or words
- ▶ Called a *line*
- ▶ Ex: On Skylake: Line size = 64 bytes

Line

- ▶ Ex if we ask for data at address 1,000: Cache will pull in data from addresses 960...1023
 - ▶ $960 = 15 * 64$; $1024 = 16 * 64$
- ▶ Idea: If we ask for data at address x , we'll probably ask for data at address $x+1$ soon

Note

- ▶ Many modern CPU's maintain separate caches for instructions & data
- ▶ Rationale: Code & data usually widely separated in RAM
- ▶ So we don't use same cache for both
- ▶ Ex: Skylake:
 - ▶ L1 = 8KB/8KB split
 - ▶ L2, L3 = *unified*

Question

- ▶ Where does data get placed in cache?
- ▶ Ex: Suppose cache is 32KB (typical size today)
 - ▶ 32KB = 512 lines in cache
- ▶ Layout of one cache entry: Several fields
 - ▶ Valid bit (is there valid data in this slot?)
 - ▶ Dirty bit (is cache same as main-memory copy?)
 - ▶ Tag bits: Identify which line of memory this holds
 - ▶ Data: The line itself
- ▶ At bootup or context switch, all “Valid” bits get set to zero
 - ▶ Also known as cache flush

Direct Mapped

- ▶ Suppose we're accessing address A
- ▶ Let $L = A \bmod 64$
- ▶ Store data to cache line L, set tag=A
 - ▶ Remember: A line is smallest quantum for cache
 - ▶ We can't load "half a line" to the cache
- ▶ This is simple & easy
- ▶ So what's the drawback?

Drawback

- ▶ Conflicts: If two addresses map to same line, they'll evict each other from cache
- ▶ If cache is 8KB (512 lines), *any* two addresses 8KB apart *must* conflict
- ▶ Power-of-two array sizes are very common in computer programs, so this problem appears frequently

Example

- ▶ Copy data between two arrays:

```
unsigned A[2048];  
unsigned B[2048];  
for(i=0;i<2048;++i){  
    A[i] = B[i];  
}
```

- ▶ Consider what happens here...

Process

- ▶ CPU fetches B[0]
 - ▶ Brings in whole line.
 - ▶ Suppose it ends up in cache slot (line) 10.
- ▶ CPU writes to A[0]. This requires bringing in that line.
 - ▶ It will end up in slot 10, so it evicts data from B
- ▶ Then CPU fetches B[1]
 - ▶ This results in a *cache miss*, so the line is brought in to slot 10
- ▶ CPU writes to A[1].
 - ▶ Another miss!
- ▶ And so on

Fully Associative

- ▶ Suppose we're accessing address A
- ▶ See if line is in cache: Check every tag entry for match with A
- ▶ If none: Scan cache for free line
 - ▶ Free = valid bit = 0
 - ▶ If no free line: Select *least recently used* line
 - ▶ Load data to that place, set tag=A
- ▶ This is the most flexible option
- ▶ So what's the drawback?

Drawback

- ▶ It's too difficult to check all 512 lines
- ▶ Every line would need to be checked in parallel
 - ▶ And: Must happen within one access operation
 - ▶ Remember: 4 cycles per access!

N-Way Associative

- ▶ Compromise between the other two strategies
- ▶ Suppose we have an 8-way set associative cache
- ▶ Each line could go in one of 8 possible places
- ▶ Ex: 8KB cache, 64 byte lines, 512 lines total
- ▶ Suppose we're accessing memory address A
- ▶ If we pull in line L, it can map to $(L \bmod 512)$, $(L+64)$, $(L+128)$, $(L+192)$, ...
 - ▶ Why 64 here?
 - ▶ $64 = \text{numberOfLines} \div \text{lineSize}$
 - ▶ $64 = 512 \div 8$
- ▶ This is the scheme most modern CPU's use

Example

- ▶ Intel Skylake:
 - ▶ L1 = 8 way associative
 - ▶ L2 = 4 way associative
 - ▶ L3 = 16 way associative

Line State

- ▶ Suppose two CPU cores accessing same area of RAM
- ▶ We'd have two copies of that line in two L1 caches
- ▶ This results in data inconsistency. What's the “true” value of that memory area?
- ▶ CPU's resolve this by using *cache protocol*

MESI


- ▶ Common protocol: MESI
- ▶ Each cache line can be in one of four states:
 - ▶ Modified
 - ▶ Exclusive
 - ▶ Shared
 - ▶ Invalid

Definition

- ▶ Invalid = This cache line doesn't contain useful data
- ▶ Shared = This line is also in some other core's cache, and it's the same in both of them
- ▶ Exclusive = This core's cache is the only one that has this data, but it's the same as in main memory
- ▶ Modified = This core's cache is the only one that has this data, and it's been modified from main memory

Transition

- Transition diagram:



Result

- ▶ Worst case: Two cores are manipulating variables that are in same line

```
int A[10];
int B[10];
int C[10];
int D[10];
void core1(){
    for(int i=0;i<10;++i)
        B[i] = A[i]+1;
}
void core2(){
    for(int i=0;i<10;++i)
        D[i] = C[i]+1;
}
```

- ▶ Consider what happens...

Result

- ▶ A, B, C, D are in same line
- ▶ Core 1 loads line for A into its cache, marks as Exclusive
- ▶ Core 1 can write to B immediately
- ▶ Suppose core 2 starts executing

Result

- ▶ Needs line. Broadcasts message asking for it
- ▶ Core 1 puts contents of line on bus, moves it to Shared
- ▶ Core 2 loads contents from bus, marks as Shared
- ▶ Core 2 writes to D[0].
 - ▶ Broadcasts message on bus to take ownership
 - ▶ Marks its copy of line as Modified
 - ▶ Core 1 marks its copy as Invalid

Result

- ▶ Now, core 1 wants to read A[1] and write B[1]
 - ▶ Broadcasts message on bus asking for it
 - ▶ Core 2 moves from Modified → Shared
 - ▶ Core 1 moves from Invalid → Shared
- ▶ Etc.

Timing

```
#define N 1000
int A[N][N], B[N][N], C[N][N];
int main(int argc, char* argv[]){
    for(int i=0;i<N;++i){
        for(int j=0;j<N;++j){
            A[i][j] = rand() & 0xff;
            B[i][j] = rand() & 0xff;
        }
    }
    clock_t st,et;  struct tms x;
    st = times(&x);
    for(int i=0;i<N;++i){
        for(int j=0;j<N;++j){
            int total=0;
            for(int k=0;k<N;++k)
                total += A[i][k] * B[k][j];
            C[i][j]=total;
        }
    }
    et = times(&x);
    //prevent optimizing it out
    cout << "C=" << C[rand()&0xff][rand()&0xff] << "\n";
    cout << et-st << "\n";
    return 0;
}
```

Result

- ▶ When I run it, I get a time of 1.03 sec
- ▶ But suppose we change this line: `total += A[i][k] * B[k][j];` to: `total += A[i][k] * B[j][k];`
 - ▶ Obviously, this wouldn't be correct matrix multiplication
 - ▶ But ignore that for now...
- ▶ The time goes down to 0.54 seconds!
- ▶ Why?

Locality

- ▶ C stores arrays in *row major order*
- ▶ So what happens when we have the line $\text{total} += A[i][k] * B[k][j]$
- ▶ Suppose $k=0$
 - ▶ Load cache with line containing $A[i][0] \dots A[i][15]$
 - ▶ Load cache with line containing $B[0][j] \dots B[0][j+15]$
 - ▶ Do multiply, store result
- ▶ So far, so good

Locality

- ▶ What happens on next iteration?
 - ▶ CPU loads $A[i][1]$ into a register
 - ▶ Fetches from cache. Very fast.
 - ▶ CPU loads $B[1][j]$ into a register
 - ▶ Not in cache!
 - ▶ A miss...Slow
- ▶ And this happens for every single array element

Locality

- ▶ But consider what happens if the statement is: $\text{total} += A[i][k] * B[j][k];$
- ▶ Again, let $k=0$
 - ▶ Load cache with line containing $A[i][0] \dots A[i][15]$
 - ▶ Load cache with line containing $B[j][0] \dots B[j][15]$
 - ▶ Do multiply, store result

Locality

- ▶ What happens on next iteration?
 - ▶ CPU loads $A[i][1]$ into a register
 - ▶ Fetches from cache. Very fast.
 - ▶ CPU loads $B[j][1]$ into a register
 - ▶ In the cache: Fast
- ▶ When do we get cache misses?

Misses

- ▶ Every time k crosses over a value of 16
 - ▶ 16, 32, 48, 64
- ▶ Note: Some CPU's have hardware prefetchers
 - ▶ If memory access pattern is regular: It starts pulling in consecutive lines
 - ▶ So we might not have any misses at all!

Prefetcher

- ▶ Prefetcher looks for access patterns and anticipates them
- ▶ Ex: Going linearly forward or backward in an array
- ▶ How to best structure accesses?

Compiler

- ▶ We can help the compiler out:
 - ▶ Avoid globals (they're stored far from locals in RAM)
 - ▶ Simple control flow (try not to nest if's, switch, loops, etc.)
 - ▶ const can help
- ▶ CPU also provides 'prefetch' instruction to allow you to specify addresses you'll need soon

Write Caching

- ▶ Caches also handle writes
 - ▶ When you write to cache line, CPU collects all writes and issues them in background
 - ▶ Called “write-back” strategy

Sources

- ▶ Intel Optimization manual
- ▶ <http://www.7-cpu.com/cpu/Skylake.html>
- ▶ <http://valgrind.org/docs/manual/cg-manual.html>
- ▶ <https://gist.github.com/jboner/2841832>
- ▶ DDR3 vs. DDR4: Four Generations, Raw Bandwidth by the Numbers.
http://www.corsair.com/en-eu/blog/2015/september/ddr3_vs_ddr4_generational
- ▶ Rajinder Gill. The Ins and Outs of Memory Addressing. <https://www.anandtech.com/show/3851/everything-you-always-wanted-to-know-about-sdram-memory-but-were-afraid-to-ask/2>
- ▶ DocMemory. <http://www.simmtester.com/page/news/showpubnews.asp?num=168>
- ▶ Thomas Schwarz, S.J. COEN 180. <http://www.cse.scu.edu/~tschwarz/coen180/LN/DRAM.html>
- ▶ Thomas Schwarz, S.J. COEN 180. http://www.cse.scu.edu/~tschwarz/COEN180_06/LN/sram.html
- ▶ Micron Corp. https://www.micron.com/products/datasheets/%7B3D323C4D-6BC7-4193-908D-E99AD746AA4E%7D?page=30#topic=concept_03583372F4E04674BE846190D4B79C89
- ▶ Intel Corp. Intel 64 and IA-32 Architectures Optimization Reference Manual.

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>