

Optimization

Motivation

- ▶ Examine some other optimization patterns
- ▶ Several C++ specific techniques

Containers

- ▶ Nearly all programs must manipulate data
- ▶ That means we need to store it in a data structure
- ▶ Choice of container can have significant impact on overall program time

Containers

- ▶ Basic container: `vector<>`
 - ▶ Indexing with `[]` or `at()` : $O(1)$
 - ▶ Append with `push_back`: *amortized constant time*: $O(1)$
 - ▶ Insert in middle/remove from middle: $O(n)$

Containers

- ▶ Linked list: `list<>`
 - ▶ No indexing possible
 - ▶ Append/insert/remove: $O(1)$, all cases

Containers

- ▶ Double ended queue: deque
- ▶ How it's stored internally: Like `vector<vector<>>`
 - ▶ Outer vector contains more vectors
 - ▶ Inner vectors contain the data items
- ▶ This leads to unique time characteristics...

Deque

- ▶ Prepending or appending is $O(1)$
- ▶ Access (with `[]`) is $O(1)$
 - ▶ But larger hidden constant than `vector`
- ▶ Insert is $O(n/2) = O(n)$

Example

- ▶ Suppose we are getting collection of data items (maybe from network)
- ▶ We want to store them in reversed order
- ▶ For benchmarking, we'll use block of memory with data items and just time the container insertion operation
- ▶ Code: [reverse.cpp](#)

Results

```
$ ./a.out 10  
Times: vector = 2 usec, deque = 3 usec  
$ ./a.out 100  
Times: vector = 4 usec, deque = 1 usec  
$ ./a.out 1000  
Times: vector = 159 usec, deque = 19 usec  
$ ./a.out 10000  
Times: vector = 4946 usec, deque = 52 usec  
$ ./a.out 100000  
Times: vector = 840208 usec, deque = 255 usec
```

Question

- ▶ What about appending?
- ▶ [append.cpp](#)

Results

```
$ ./a.out 10  
Times: vector = 1 usec, deque = 4 usec  
$ ./a.out 100  
Times: vector = 1 usec, deque = 3 usec  
$ ./a.out 1000  
Times: vector = 13 usec, deque = 23 usec  
$ ./a.out 10000  
Times: vector = 40 usec, deque = 49 usec  
$ ./a.out 100000  
Times: vector = 361 usec, deque = 404 usec  
$ ./a.out 1000000  
Times: vector = 3605 usec, deque = 4254 usec
```

Insert

- ▶ What about inserting in the middle?
- ▶ [middle.cpp](#)

Results

```
$ ./a.out 10
```

```
Times: vector = 3 usec, deque = 2 usec
```

```
$ ./a.out 100
```

```
Times: vector = 4 usec, deque = 10 usec
```

```
$ ./a.out 1000
```

```
Times: vector = 102 usec, deque = 376 usec
```

```
$ ./a.out 10000
```

```
Times: vector = 1536 usec, deque = 4856 usec
```

```
$ ./a.out 100000
```

```
Times: vector = 366951 usec, deque = 533487 usec
```

Parameters

- ▶ Accidental copying of function parameters can consume considerable time

```
void foo( BigThing b ){  
    ...  
}
```

Solution?

- ▶ We can avoid overhead of copies by using references:

```
void foo( BigThing& b ){  
    ...  
}
```

- ▶ But what if someone accidentally tries to copy object with =?

C++ 11

- ▶ C++ 11 provides *deleted* functions

```
class Foo{  
    public:  
    void operator=(const Foo& x) = delete;  
    Foo(const Foo&) = delete;  
};
```


Benefit

- ▶ This catches several problems:

- ▶ Pass by value

- ```
void foo(BigThing b){ ... }
```

- ▶ Use of =

- ```
BigThing b2 = b;
```

- ▶ Use of copy constructor

- ```
BigThing b2(b);
```

- ▶ Return by value

- ```
BigThing foo(){ ... }
```

Loop Tests

- ▶ Consider this code:

```
char x[512];  
...put data in x...  
for(auto i=0;i<strlen(x);i++){  
    ...examine x[i]...  
}
```

- ▶ What's the problem?

Problem

- ▶ `strlen` is $O(n)$
- ▶ Called on every loop iteration
- ▶ If length of `x` not changed in loop, this is wasteful

Better

- Compute length once, cache it

```
char x[512];  
...put data in x...  
auto len = strlen(x);  
for(suto i=0;i<len;i++){  
    ...examine x[i]...  
}
```

Or

- ▶ Do the count the other way:

```
char x[512];  
...put data in x...  
for(auto i=strlen(x)-1;i>=0;i--){  
    ...examine x[i]...  
}
```

- ▶ This code is broken. Do you know why?

Problem

- ▶ Consider: `for(auto i=strlen(x);i>=0;i--)`
- ▶ What type is `i`?

Type

- ▶ `strlen` defined to return `size_t`
- ▶ `size_t` is typically either `uint32_t` or `uint64_t`
- ▶ So test is like: `for(unsigned i=strlen(x)-1;i>=0;i--)`
- ▶ When does the termination condition get triggered?

Problem

- ▶ Never!
- ▶ On last iteration of loop: $i=0$
- ▶ Then $i--$ occurs
- ▶ i wraps to `INT_MAX`
- ▶ And loop keeps going!

Solution

- ▶ Need to use signed type here
- ▶ But: If string is very long (2GB+), 'i' will be initialized to negative value

Note

- ▶ Same problem can occur in other contexts
- ▶ Ex:

```
vector<Foo> v;  
...put stuff in v...  
for(auto i=v.size()-1;i>=0;i--){  
    ...  
}
```

Idea

- ▶ Since loops run many times, they are suspect for code hot spots
- ▶ One rule of thumb: Avoid function calls in loop
- ▶ Some are not obvious

Question

- ▶ How many function calls are made?

```
for(int i=0;i<100;++i){  
    Foo f;  
}
```

Question

- ▶ How many function calls are made?

```
for(int i=0;i<100;++i){  
    Foo f;  
}
```

- ▶ 200!
 - ▶ Constructor runs at top of loop
 - ▶ Destructor runs at bottom

Example

- ▶ Example (based on one in Guntheroth):
- ▶ This is suboptimal: Lots of constructor/destructor calls:

```
for(...){  
    string x("something");  
    ...code that changes x...  
}
```

Improved

- ▶ We can improve this by taking advantage of the string class's abilities:

```
string x;  
for(...){  
    x.clear()  
    x+="...";  
    ...code that changes x...  
}
```

Loops

- ▶ Be aware of overloaded operators
 - ▶ $+$, $-$, $*$, $/$, etc.
 - ▶ Operations can be more costly than they look
- ▶ Ex:
mat4 A = ...;
mat4 B = ...;
mat4 C = A*B;
- ▶ The matrix-matrix multiply is much more expensive than a single scalar multiply

Organization

- ▶ Instead of calling a function in a loop, it can be better to perform a loop in a function
- ▶ Ex: Suboptimal:

```
vector<...> foo;  
for(i=0;i<foo.size();i++){  
    func(foo[i]);  
}
```

Better:

```
funcAll(foo);  
void funcAll(vector<...>& foo ){  
    for(i=0;i<foo.size();i++){  
        ...process foo[i]...  
    }  
}
```

Explanation

- ▶ Why is first one potentially slower?
 - ▶ We have function call overhead on every loop iteration
 - ▶ Second form has only one function call


Classes

► Consider inheritance:

```
class Base{  
    int x,y;  
};  
class Der : public Base {  
    int z;  
    virtual void foo(){  
    }  
};
```

Vtable

- ▶ If class contains virtual functions, it must have a vtable
- ▶ Organization of Base and Der:



Example

- ▶ Suppose we have:

Der* d = ...;

d->foo();

- ▶ System must get value in d (the address), add sizeof(Base) to it, then do indirect function call on result
 - ▶ This can be slower


Better

```
class Base{
    int x,y;
    virtual ~Base(){}
};
class Der : public Base {
    int z;
    virtual void foo(){
    }
};
```

- ▶ Any function that's going to be the base of inheritance chain should have virtual destructor anyway

Result

- Organization of memory:



- Now, vtable is at start of object
 - Saves an addition upon function call

Multiple Inheritance

- ▶ C++ also allows multiple inheritance:

```
class Foo : public Bar, public Baz {
```

```
    ...
```

```
};
```

- ▶ Problem: If Bar and Baz both have virtual functions, Foo has multiple vtables
 - ▶ This results in the same issue with needing to add an offset to pointer location to get function addresses
 - ▶ So multiple inheritance can be (a bit) slower to use

Inlining

- ▶ In C++, if function body inside class declaration, it's usually inlined
 - ▶ Saves function call/return time
 - ▶ Might help optimize register usage
 - ▶ But: Virtual functions normally cannot be inlined
- ▶ So: Virtual functions can have additional cost here too
- ▶ But: Series of switch/if-else can be slower than virtual function call

Sources

- ▶ Herb Sutter. GotW: Compilation Firewalls.
https://herbsutter.com/gotw/_100/
- ▶ <https://stackoverflow.com/questions/22306949/does-deque-provide-o1-complexity-when-inserting-on-top>
- ▶ <https://stackoverflow.com/questions/7572529/complexity-of-stl-dequeinsert>
- ▶ <https://stackoverflow.com/questions/6292332/what-really-is-a-deque-in-stl>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>