

Profiling 1

Quote

*Measure what is measurable, and make measurable what is not so.
—Galileo Galilei, as reported by Kurt Guntheroth.*

Motivation

- ▶ We have a beautiful raytracing program
- ▶ But we'd like it to finish its work more quickly
- ▶ Is our only solution to spend more money on a faster computer?

Tools

- ▶ We have three tools we can use
 - ▶ Profiler: Compiler will tell you where the “hot spots” are
 - ▶ Instrumentation: We can add code to measure the hot spots
 - ▶ Notes: Important to keep track of our observations
 - ▶ Either pencil and paper or else in a text file

Optimizing

- ▶ Optimization is much like scientific experiments:
 - ▶ Prediction: Predict where slow spots are
 - ▶ Important to commit it to writing
 - ▶ Memory is unreliable; writing it down helps to focus, keep records
 - ▶ So we don't go in circles!

Optimizing

- ▶ Optimization is much like scientific experiments:
 - ▶ Prediction: Predict where slow spots are
 - ▶ Record changes that have been made to code
 - ▶ Either use version control comments
 - ▶ Or else write down on paper or in text file

Optimizing

- ▶ Optimization is much like scientific experiments:
 - ▶ Prediction: Predict where slow spots are
 - ▶ Record changes that have been made to code
 - ▶ Measure

Optimizing

- ▶ Optimization is much like scientific experiments:
 - ▶ Prediction: Predict where slow spots are
 - ▶ Record changes that have been made to code
 - ▶ Measure
 - ▶ Record results (in writing!)

Rule of Thumb

- ▶ Pareto's law: 80/20 rule
 - ▶ Sometimes quoted as 90/10 rule

Rule

- ▶ Amdahl's law: $\text{Speedup} = \frac{T_{\text{slow}}}{T_{\text{fast}}}$
- ▶ Normally, we think of this in context of parallelization
- ▶ But it also applies for optimization

Example

- ▶ Suppose one piece of code accounts for 80% of runtime
- ▶ And we can make it take three-quarters the time it previously took
- ▶ Overall program speedup: $\frac{1}{0.8*0.75+0.2*1} = 1.25$

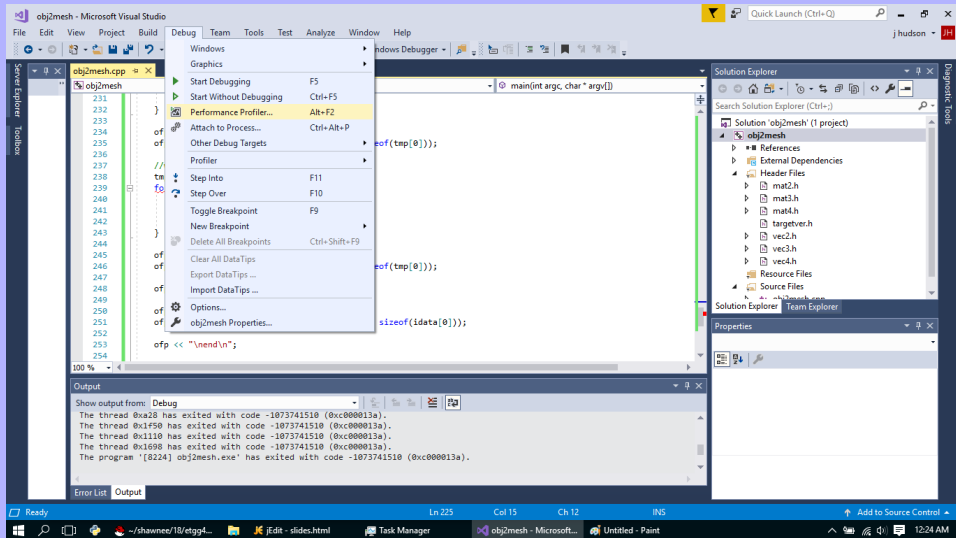
Conversely

- ▶ Suppose a piece of code accounts for only 30% of runtime
- ▶ We make it take 75% of the time it previously took
- ▶ Overall program speedup: $\frac{1}{0.3*0.75+0.7*1} = 1.08$
 - ▶ Overall program is only 8% faster
- ▶ So we want to concentrate our efforts where they will have larger payoff

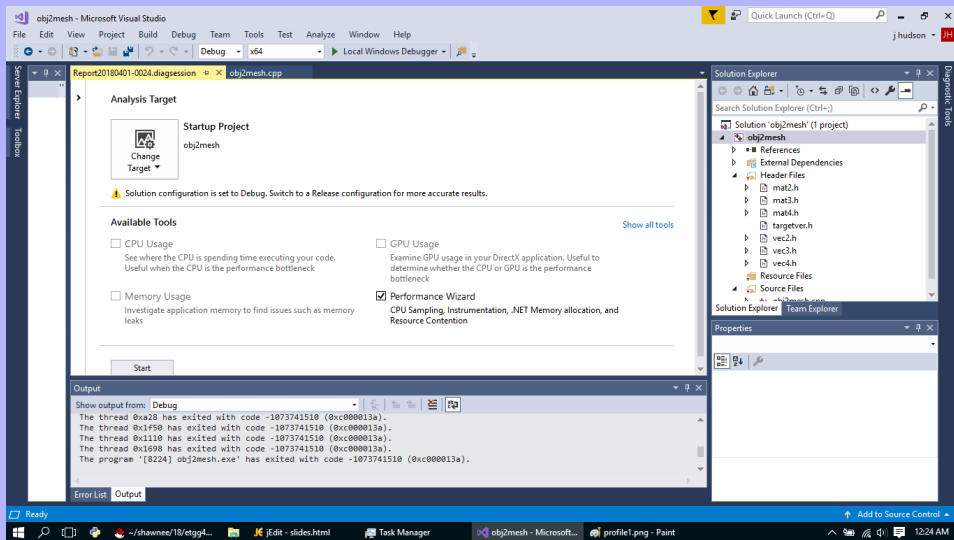
Profiling

- ▶ Most compilers/runtime systems today have profiling systems built-in
- ▶ We'll look at C/C++ (Visual Studio & Valgrind)


VS



VS



Performance Wizard -- Page 1



Specify the profiling method

Profiling your application can help diagnose performance problems and identify the most common expensive methods in your application. To begin, choose a profiling method from the options below.

What method of profiling would you like to use?


- ☒ **CPU sampling (recommended)**
Monitor CPU-bound applications with low overhead
- ☐ **Instrumentation**
Measure function call counts and timing
- ☐ **.NET memory allocation**
Track managed memory allocation
- ☐ **Resource contention data (concurrency)**
Detect threads waiting for other threads

[Read more about profiling methods](#)

< Previous Next > Finish Cancel

VS

Performance Wizard -- Page 2

 Choose the application to profile using the sampling method

Choose the application that you want to profile (.EXE, .DLL, Website)

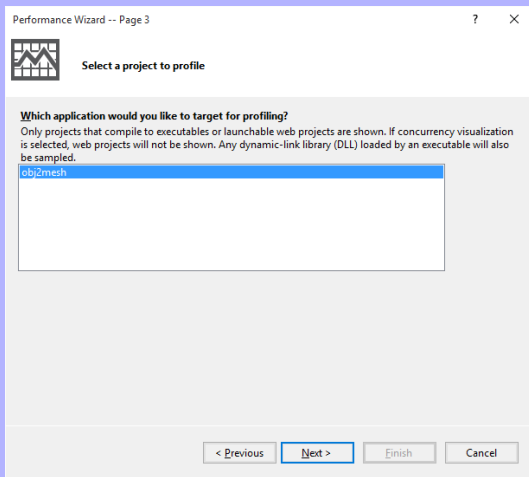
☒ One or more available projects

☐ An executable (.EXE file)

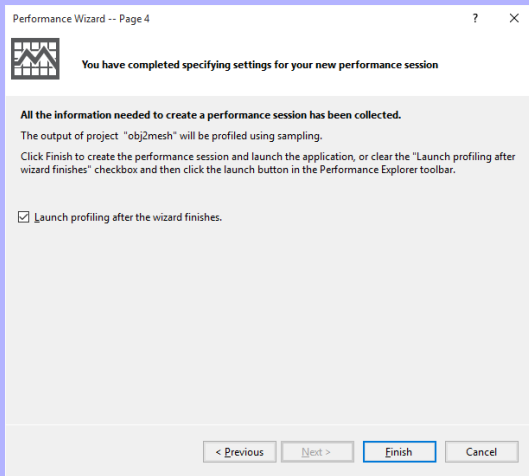
☐ An ASP.NET application

< Previous **Next >** Finish Cancel

VS



VS



Methods

- ▶ Two methods for C code:
 - ▶ Sampling
 - ▶ Instrumentation

Sampling

- ▶ Sampling = CPU is periodically interrupted by profiler
 - ▶ Profiler examines eip/rip register, sees which function is executing
 - ▶ Records this information
 - ▶ Default: MS uses once per 10M cycles
- ▶ Advantage: Less intrusive: Executing code is not modified at all
- ▶ Disadvantage: Less exact
 - ▶ The more samples, the more exact
 - ▶ But more overhead too
- ▶ Disadvantage: Sometimes it gives no data!

Instrumentation

- ▶ Every function executes extra code in prologue/epilogue
- ▶ Advantage: More exact: Every function call is caught
- ▶ Disadvantage: More overhead

Example

- ▶ (Do in class: Example executions)
- ▶ Sampling: Exclusive vs. Inclusive counts
 - ▶ On each interrupt, we have one function ip is inside of and other function(s) on call stack
 - ▶ Function we're inside of: Gets *exclusive* count incremented
 - ▶ Other functions: Get *inclusive* count incremented

Linux

- ▶ `gcc -g -O3 main.cpp -o main`
- ▶ `valgrind --tool=callgrind ./main scene1.txt`
- ▶ `callgrind_annotate --show-percs=yes --auto=yes callgrind.out.*`

Example

- ▶ I ran my raytracing code under valgrind on Linux on a scene with 3 spheres and one esh and found that ~92% of the time was in traceTriangles

```
12,199,146,303 (69.07%)  math3d.h:traceTriangles(Scene&, vec3Base<float>&, vec3Base<float>&, vec3Base<float>&,
                        vec3Base<float>&, vec3Base<f
loat>&, float&, bool&)
4,044,274,321 (22.90%)  traceRay.h:traceTriangles(Scene&, vec3Base<float>&, vec3Base<float>&, vec3Base<float
                        >&, vec3Base<float>&, vec3Base<float>&, float&, bool&) [main-o3]
1,208,483,840 ( 6.84%)  /usr/include/c++/9/bits/stl_vector.h:traceTriangles(Scene&, vec3Base<float>&,
                        vec3Base<float>&, vec3Base<float>&, vec3Base<float>&, float&, bool&)
63,006,219 ( 0.36%)    ???:0x00000000000027e90 [/usr/lib64/libpng16.so.16.36.0]
```

- ▶ Notice that (due to optimization) profiler thinks traceTriangles is spread out between several source code files

Profiling

► Let's drill down into the functions...

```
void traceSpheres(Scene& scene, vec3& s, vec3& v, vec3& ip, vec3& N, vec3& color, float& closestT, bool& inter)
{
    unsigned ci=(unsigned)-1;
    for(unsigned i=0;i<scene.spheres.size();++i){
        auto q = s - scene.spheres[i].c;
        auto A = dot(v,v);
        auto B = 2*dot(q,v);
        auto C = dot(q,q)-scene.spheres[i].r*scene.spheres[i].r;
        auto disc = B*B-4*A*C;
        if( disc < 0 )
            continue;
        disc = sqrt(disc);
        auto denom = 1.0 / (2.0*A);
        auto t1 = (-B + disc) * denom;
        auto t2 = (-B - disc) * denom;
        float t;
        if( t1 < 0 && t2 < 0 ){
            continue;
        } else if( t1 < 0 )
            t = t2;
        else if( t2 < 0 )
            t = t1;
        else
            t = ( t1<t2 ) ? t1:t2 ;

        if( t < closestT ){
            closestT = t;
            ci = i;
        }
    }
    if( ci == (unsigned)-1 ){
        inter=false;
        return;
    }
    ip = s + closestT * v;
    N = normalize(ip - scene.spheres[ci].c);
    color = scene.spheres[ci].color;
    inter=true;
}
```

1,572,864 (0.01%)
262,144 (0.00%)
5,242,880 (0.03%)
786,432 (0.00%)
1,572,864 (0.01%)
2,359,296 (0.01%)
1,572,864 (0.01%)
60,024 (0.00%)
80,032 (0.00%)
60,024 (0.00%)
604,320 (0.00%)
60,024 (0.00%)
40,016 (0.00%)
40,016 (0.00%)
60,024 (0.00%)
19,980 (0.00%)
19,980 (0.00%)
524,288 (0.00%)
484,328 (0.00%)
79,920 (0.00%)
59,940 (0.00%)
79,920 (0.00%)
39,960 (0.00%)

Analysis

- ▶ We see that traceSpheres doesn't account for much overall runtime
- ▶ But: This input file had three spheres and 768 triangles
- ▶ We should run program on more input files to get better sense of where bottlenecks are

Profile

► What about traceTriangles?

```
void traceTriangles(Scene& scene, vec3& s, vec3& v, vec3& ipC, vec3& N,
                    vec3& color, float& closestT, bool& inter)
{
    2,621,440 ( 0.01%) for(auto& M : scene.meshes ){
    1,572,864 ( 0.01%)     for(unsigned i=0;i<M.triangles.size();++i){
    807,141,376 ( 4.57%)         Triangle& T = M.triangles[i];
                                float denom = dot(T.N,v);
                                if( denom == 0.0 )
                                continue;
    805,306,368 ( 4.56%)         float numer = -(T.D + dot(T.N,s) );
                                float t = numer/denom;
                                if( t < 0 )
                                continue;
                                auto ip = s + t * v;
                                float A=0.0;
                                for(unsigned j=0;j<3;++j)
    472,522,149 ( 2.68%)                 A += length( cross(T.e[j], ip-T.p[j] ) );
    472,522,149 ( 2.68%)         A *= T.oneOverTwiceArea;
    473,046,437 ( 2.68%)         if( A > 1.001 )
                                continue;
                                if( t < closestT ){
    163,260 ( 0.00%)                 ipC = ip;
    82,719 ( 0.00%)                 N = T.N;
    110,292 ( 0.00%)                 color = M.color;
    110,292 ( 0.00%)                 closestT = t;
    27,573 ( 0.00%)                 inter = true;
    55,146 ( 0.00%)             }
                                }
    }
    2,097,152 ( 0.01%) }
```

Analysis

- ▶ Notice that this function doesn't seem to be very significant time-wise either!
- ▶ But...

Profile

► Take a look at traceRay:

```
.      bool traceRay(Scene& scene, vec3& s, vec3& v, vec3& ip, vec3& N, vec3&
      color)
.      {
262,144 ( 0.00%)      float closestT = 1E99;
.                      bool inter;
2,359,296 ( 0.01%)      traceSpheres(scene,s,v,ip,N,color,closestT,inter);
35,782,500 ( 0.20%) => traceRay.h:traceSpheres(Scene&, vec3Base<float>&, vec3Base<float>&,
      vec3Base<float>&, vec3Base<float>&, vec3Base<float>&, float&, bool&) (262,144x)
2,359,296 ( 0.01%)      traceTriangles(scene,s,v,ip,N,color,closestT,inter);
17,452,166,608 (98.82%) => traceRay.h:traceTriangles(Scene&, vec3Base<float>&, vec3Base<float>
      &, vec3Base<float>&, vec3Base<float>&, vec3Base<float>&, float&, bool&) (262,144x)
.                      return inter;
.      }
```

Analysis

- ▶ Notice that this says that `traceTriangles` is responsible for >98% of the program's execution time but `traceTriangles` doesn't seem to be anywhere near that costly
- ▶ What's going on????

Explanation

- ▶ This report is showing only *exclusive* counts
- ▶ We need to look through the rest of the report for high-cost operations
- ▶ When I did so, this is what I found...

Profile

```
.      float dot(const vec3& v, const vec3& w){
3,231,002,556 (18.29%)      return v.x*w.x + v.y*w.y + v.z*w.z;
.      }

1,422,912,839 ( 8.06%)      MyType operator-(const MyType& other) const {      op3(-);      }

.      vec3 cross(const vec3& v, const vec3& w){
1,417,568,757 ( 8.03%)      return vec3(
472,522,919 ( 2.68%)          v.y*w.z - w.y*v.z,
945,045,068 ( 5.35%)          w.x*v.z - v.x*w.z,
1,417,567,217 ( 8.03%)          v.x*w.y - w.x*v.y
.      );
.      }

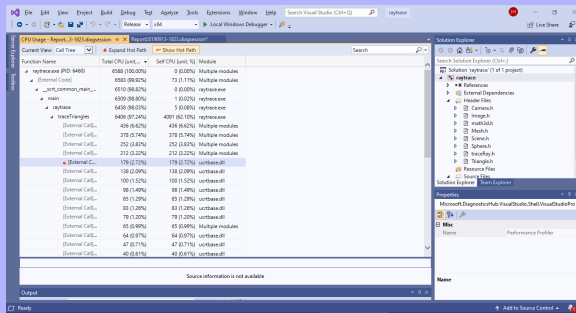
.      float magnitudeSq() const {
.          float total;
.          total = this->x*this->x;
945,230,236 ( 5.35%)      total += this->y*this->y;
945,230,236 ( 5.35%)      total += this->z*this->z;
.          return total;
.      }
```

VS

- ▶ Let's see what Visual Studio has to say...

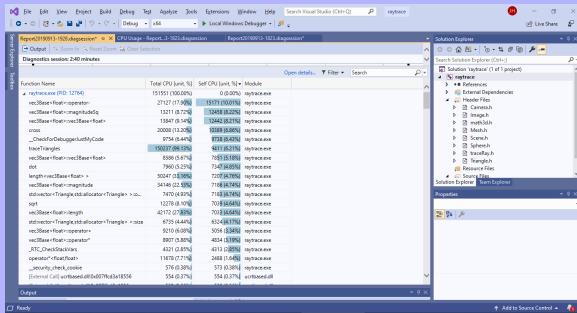
VS

- ▶ Using Release mode, we may not get very useful results



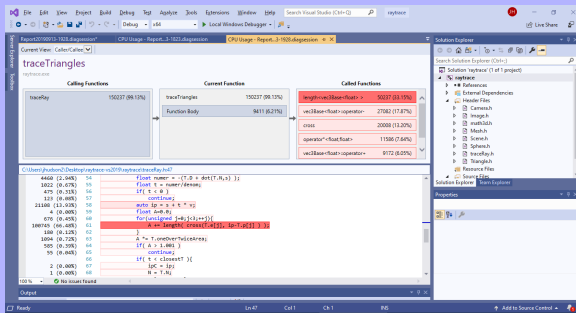
VS

- VS warns against using Debug mode, but it gives us more information



VS

- ▶ We can zoom in on a single function (double click its name)



Analysis

- ▶ We can see that no single statement is responsible for a large chunk of the execution time
 - ▶ If it was, that would be an obvious place to start our optimization efforts

Analysis

- ▶ Much of our time is being spent in fundamental operations like dot, cross, subtraction
- ▶ Not really clear how we could speed these up (other than by doing fewer of them)

Speeding Up

- ▶ Can we make the program faster, though?
- ▶ Yes!
- ▶ A few strategies that we can apply in these cases

Technique

- ▶ Don't recompute data in a loop if we don't need to
- ▶ Look carefully at sphere intersection routines. Is there anything being repeatedly computed that need not be?
 - ▶ Go back a few slides, have a look...

Technique

- ▶ Look more closely:

```
auto q = s - scene.spheres[i].c;  
auto A = dot(v,v);  
auto B = 2*dot(q,v);  
auto C = dot(q,q)-scene.spheres[i].r*scene.spheres[i].r;
```

- ▶ Observe: A depends only on v
 - ▶ Which is constant for all iterations of loop
- ▶ Observe: C depends on r^2 which is constant for each sphere

Code

- I rewrote the code like so:

```
auto A = dot(v,v);
auto fourA = 4.0*A;
auto denom = 2.0 / (fourA);
unsigned ci=(unsigned)-1;
for(unsigned i=0;i<scene.spheres.size();++i){
    auto q = s - scene.spheres[i].c;
    auto B = 2*dot(q,v);
    auto C = dot(q,q)-scene.spheres[i].r2;
    auto disc = B*B-fourA*C;
    if( disc < 0 )
        continue;
    disc = sqrt(disc);
    auto t1 = (-B + disc) * denom;
    auto t2 = (-B - disc) * denom;
    ...rest is the same...
```

Results

- ▶ I ran the code on an input file with 1000 spheres
 - ▶ Original code: 1.145 sec
 - ▶ New code: 1.152 sec
- ▶ Hmm...

Explanation

- ▶ Modern compilers are very good at hoisting loop invariants
- ▶ So we don't really see any significant benefit from doing the work ourselves

Technique

- ▶ Second technique: Examine the algorithm we're using and make it better
 - ▶ Advantage: This can give significant speedups
 - ▶ Disadvantage: No general way to approach this
 - ▶ You just need to study your code carefully and *reason* it through
- ▶ Let's see some examples

Example

- ▶ Recall our scheme to find ray-sphere intersections:
 - ▶ Ray start, direction are \vec{s} & \vec{v} ; sphere center, radius are \vec{c} & r ; $\vec{q} = \vec{s} - \vec{c}$
- ▶ $t = \frac{-B \pm \sqrt{B^2 - 4AC}}{2A}$
 - ▶ $A = \vec{v} \cdot \vec{v}$
 - ▶ $B = 2(\vec{q} \cdot \vec{v})$
 - ▶ $C = \vec{q} \cdot \vec{q} - r^2$
 - ▶ If discriminant < 0 : No intersection
 - ▶ Otherwise, take smallest value of t that is nonnegative, compare it to the closest t value found so far
- ▶ This can be made better!

Observe

- ▶ We defined $B=2(\vec{q} \cdot \vec{v})$
- ▶ But what if we actually plug B into the quadratic formula?

$$\text{▶ } t = \frac{-(2(\vec{q} \cdot \vec{v})) \pm \sqrt{(2(\vec{q} \cdot \vec{v}))^2 - 4AC}}{2A}$$

$$\text{▶ } t = \frac{-2(\vec{q} \cdot \vec{v}) \pm \sqrt{4(\vec{q} \cdot \vec{v})^2 - 4AC}}{2A} = \frac{-2(\vec{q} \cdot \vec{v}) \pm \sqrt{4((\vec{q} \cdot \vec{v})^2 - AC)}}{2A}$$

$$\text{▶ } t = \frac{-2(\vec{q} \cdot \vec{v}) \pm 2\sqrt{(\vec{q} \cdot \vec{v})^2 - AC}}{2A} = \frac{-(\vec{q} \cdot \vec{v}) \pm \sqrt{(\vec{q} \cdot \vec{v})^2 - AC}}{A}$$

Observe


- ▶ We defined $A = \vec{v} \cdot \vec{v}$
 - ▶ If the vector is unit length, we know $A=1$
 - ▶ So: Normalize the ray direction and then simplify further:
 - ▶ $t = -(\vec{q} \cdot \vec{v}) \pm \sqrt{(\vec{q} \cdot \vec{v})^2 - C}$

Branches

- ▶ We always take the smallest (nonnegative) t value
- ▶ Right now, this involves branchy code
 - ▶ We saw how that can result in slower CPU performance due to branch misprediction
- ▶ We'd like to reduce or eliminate branches

Branches

- ▶ Let a and b represent the two values of t where the ray intersects the sphere
- ▶ Let $a = -(\vec{q} \cdot \vec{v}) - \sqrt{(\vec{q} \cdot \vec{v})^2 - C}$, $b = -(\vec{q} \cdot \vec{v}) + \sqrt{(\vec{q} \cdot \vec{v})^2 - C}$
- ▶ Clearly, $b \geq a$
- ▶ We can have two possibilities:



Upshot

- ▶ If we know that the ray always starts outside the sphere, we can just use $t = -(\vec{q} \cdot \vec{v}) - \sqrt{(\vec{q} \cdot \vec{v})^2 - C}$
- ▶ Note that we must still check for $t \geq 0$

Code

```
unsigned ci=(unsigned)-1;
v = normalize(v);
for(unsigned i=0;i<scene.spheres.size();++i){
    auto q = s - scene.spheres[i].c;
    auto C = dot(q,q)-scene.spheres[i].r2;
    auto qDotv=dot(q,v);
    auto disc = qDotv*qDotv-C;
    if( disc < 0 )
        continue;
    float t = -qDotv-sqrt(disc);
    if( t>0 && t < closestT ){
        closestT = t;
        ci = i;
    }
}
```

Results

- ▶ Does this give us a significant improvement?
- ▶ I tested on a scene with 1000 spheres, 512x512 resolution
 - ▶ Original code: 1.145 sec
 - ▶ New code: 0.841 sec
 - ▶ Speedup: 1.36

Triangles

- ▶ What about the triangle code?
 - ▶ If we know that our platform implements IEEE arithmetic: Anything divided by zero is ∞ , so we need not check for division by zero

Triangles

- ▶ Our intersection test involves a call to `length()`, which means taking a square root
 - ▶ Square root is typically pretty expensive
- ▶ We can use an alternate formulation: The *scalar triple product*:
 $(\vec{u} \times \vec{v}) \cdot \vec{w}$
 - ▶ Suppose we have a tetrahedron with vertices a, b, c, d
 - ▶ We can compute the volume of this tetrahedron by computing
 $((\vec{b} - \vec{a}) \times (\vec{c} - \vec{a})) \cdot (\vec{d} - \vec{a})$

Point In Triangle

- ▶ Given a triangle with vertices p_0, p_1, p_2 and potential intersection point $\vec{ip} = \vec{s} + t\vec{v}$:
 - ▶ Compute volume of tetrahedron with vertices $\vec{s}, \vec{ip}, p_0, p_1$
 - ▶ Compute volume of tetrahedron with vertices $\vec{s}, \vec{ip}, p_0, p_2$
 - ▶ Compute volume of tetrahedron with vertices $\vec{s}, \vec{ip}, p_1, p_2$
- ▶ If any volume is negative, intersection point is outside triangle; else, it's inside

Implementation

- ▶ Helper function:

```
bool scalarTripleIsNegative(const vec3& a, const vec3& b, const  
    vec3& c){  
    return dot(cross(a,b),c)<0.0f;  
}
```

Intersection

► The intersection testing code:

```
for(auto& M : scene.meshes ){
    for(unsigned i=0;i<M.triangles.size();++i){
        Triangle& T = M.triangles[i];
        float denom = dot(T.N,v);    //if zero: t=infinity
        float numer = -(T.D + dot(T.N,s) );
        float t = numer/denom;
        if( t < 0 || t >= closestT )
            continue;
        auto vv = t*v; //(s + t*v) - s
        auto v0 = T.p[0]-s;
        auto v1 = T.p[1]-s;
        auto v2 = T.p[2]-s;
        if( scalarTripleIsNegative( vv, v0,v2 ) ||
            scalarTripleIsNegative( vv, v1,v0 ) ||
            scalarTripleIsNegative( vv, v2,v1 ) ){
            continue;
        }
        auto ip = s + vv;
        ipC = ip;
        N = T.N;
        color = M.color;
        closestT = t;
        inter = true;
    }
}
```

Results

- ▶ For scene2.txt
- ▶ Original code:
 - ▶ 15.60 seconds (Linux, Core i7-5500, 2.4GHz, gcc 9.2.1, 64 bit, -O3 -march=native)
 - ▶ 24.88 seconds (Windows 10, Core i7-6500U, 2.5GHz, VS 2019, 64 bit, Release, Favor speed, use AVX instruction set)
 - ▶ 23.53 seconds (Windows, no debugger active)
- ▶ New code:
 - ▶ 8.08 seconds (Linux),
 - ▶ 10.89 seconds (Windows, no debugger active)
- ▶ Speedup: 1.93 (Linux), 2.16 (Windows) (Twice as fast!)

Principle

- ▶ Understanding the underlying mathematics/algorithms can enable us to obtain significant benefits

Observe

- ▶ The values v_0 , v_1 , v_2 are constant for a particular ray
- ▶ Maybe we'd see a speedup if we precomputed and saved these values for the triangles...

Code

```
for(auto& M : scene.meshes ){
    for(unsigned i=0;i<M.triangles.size();++i){
        Triangle& T = M.triangles[i];
        float denom = dot(T.N,v);    //if zero: t=infinity
        float numer = -(T.D + dot(T.N,s) );
        float t = numer/denom;
        if( t < 0  || t >= closestT )
            continue;
        auto vv = t*v; //(s + t*v) - s
        if( scalarTripleIsNegative( vv, T.pMinusS[0],T.pMinusS[2]) ||
            scalarTripleIsNegative( vv, T.pMinusS[1],T.pMinusS[0]) ||
            scalarTripleIsNegative( vv, T.pMinusS[2],T.pMinusS[1]) ){
            continue;
        }
        auto ip = s + vv;
        ipC = ip;
        N = T.N;
        color = M.color;
        closestT = t;
        inter = true;
    }
}
```

Note

► We also must tweak traceRay:

```
bool traceRay(Scene& scene, vec3& s, vec3& v, vec3& ip, vec3& N, vec3& color)
{
    float closestT = 1E99;
    bool inter;
    //-----new code
    static bool didpMinusS=false;
    if(!didpMinusS ){
        didpMinusS=true;
        for(auto& M : scene.meshes ){
            for(auto& T : M.triangles ){
                for(int i=0;i<3;++i){
                    T.pMinusS[i] = T.p[i] - s;
                }
            }
        }
    }
    //-----new code
    traceSpheres(scene,s,v,ip,N,color,closestT,inter);
    traceTriangles(scene,s,v,ip,N,color,closestT,inter);
    return inter;
}
```


Results

- ▶ Original code: 15.60 seconds
- ▶ Point-in-triangle using tetrahedron instead of barycentric: 8.08 seconds
- ▶ Tetrahedron + precomputation: 7.56 seconds

Notice

- ▶ We have this code at the top of our function:

```
for(unsigned i=0;i<M.triangles.size();++i){  
    Triangle& T = M.triangles[i];  
    ...  
}
```

- ▶ Question: How much time does this take?

Size

- ▶ I looked up the implementation of `size()` in my compiler (g++, libstdc++ 9.1.1)

```
size_type  
size() const _GLIBCXX_NOEXCEPT  
{ return size_type(this->_M_impl._M_finish - this->_M_impl.  
    _M_start); }
```

- ▶ That's ugly!
- ▶ But strip out the extraneous stuff and we see: It does an integer subtraction when called

[]

- ▶ What about the [] operator?

```
reference  
operator[](size_type __n) _GLIBCXX_NOEXCEPT  
{  
    __glibcxx_requires_subscript(__n);  
    return *(this->_M_impl._M_start + __n);  
}
```

- ▶ This involves an addition (and a hidden multiplication, since `_M_start` is a pointer)

Question

- ▶ What if we used the C++ 11 for-each iterator?

```
for(auto& T : M.triangles ){  
    ...  
}
```

Results

- ▶ Original code: 15.60 seconds
- ▶ Point-in-triangle using tetrahedron instead of barycentric: 8.08 seconds
- ▶ Tetrahedron + precomputation: 7.56 seconds
- ▶ Tetrahedron + precomputation + for-each: 7.35 seconds

Sources

- ▶ Kurt Guntheroth. *Optimized C++*. O'Reilly Media.
- ▶ Christer Ericson. *Real Time Collision Detection*. CRC Press.

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>