

# Compute Shaders

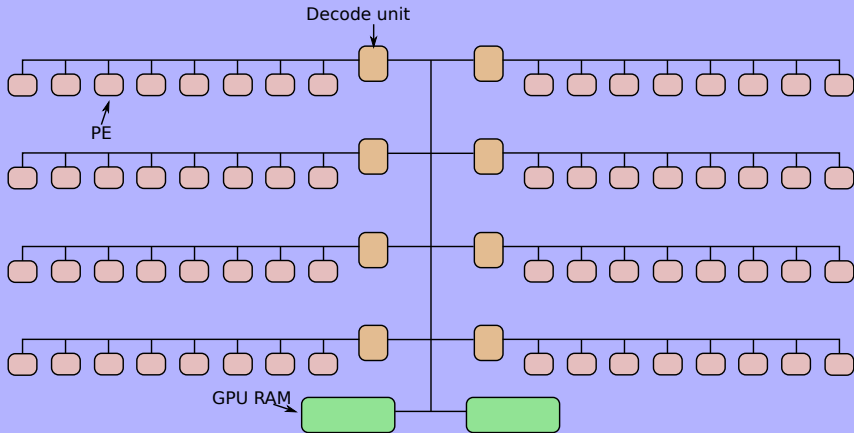
# Motivation

- ▶ You have a supercomputer inside of your computer: The GPU
  - ▶ GPU: 16,000,000,000,000 ops/sec (or more!) [nVidia Titan RTX GeForce 20 series:  
[https://en.wikipedia.org/wiki/List\\_of\\_Nvidia\\_graphics\\_processing\\_units](https://en.wikipedia.org/wiki/List_of_Nvidia_graphics_processing_units)]
  - ▶ CPU: 74,000,000,000 ops/sec (Ryzen 7, 8 core)
- ▶ We want to take advantage of this
- ▶ But...Why use CPU's at all if GPU's so fast?

# Architecture

- ▶ Consider how instructions processed on typical CPU's: F,D,O,X,S
  - ▶ Instruction is fetched, then decoded, then operands are fetched, then it's executed, and finally, the results are stored
- ▶ Every clock tick, we advance the processing of an instruction down the pipeline
- ▶ So each instruction takes five cycles in this model
- ▶ Now consider GPU architecture...

# GPU Architecture



# Architecture

- ▶ To make circuitry easier to fabricate, a set of *processor elements* are grouped together
- ▶ What's a PE?
  - ▶ Circuits to do simple operations: Adder, multiplier, divider, bitwise ops
- ▶ PE can do arithmetic, but doesn't make decisions of *what* to do
  - ▶ Decode unit does that
- ▶ So all PE's of a group must do *same* instruction but with different data
  - ▶ SIMD: Like SSE, but writ large

## Quote

- ▶ “If you were plowing a field, which would you rather use: Two strong oxen or 1024 chickens?” [Seymour Cray, early supercomputer pioneer:  
[https://www.azquotes.com/author/23611-Seymour\\_Cray](https://www.azquotes.com/author/23611-Seymour_Cray)]
- ▶ CPU = Oxen, GPU PE = Chicken

## Upshot

- ▶ GPU works best if we have very regular operations, doing many data items at once
- ▶ Not so good if we have lots of independent operations, many decisions (“branchy”) code
  - ▶ Might have to execute *both sides* of the branch
  - ▶ Discard one set of results

# Setup

- ▶ Several different platforms available:
  - ▶ OpenCL
  - ▶ CUDA
  - ▶ DirectCompute/DX 12 Compute Shaders
  - ▶ OpenGL/Vulkan Compute Shaders
- ▶ We'll use GL compute shaders here



## Prerequisites

- ▶ Download:
  - ▶ SDL 2: <https://www.libsdl.org/download-2.0.php>
  - ▶ Framework: [framework.zip](#)
- ▶ Unzip these somewhere

## Prerequisites

- ▶ Create a project file in your favorite IDE
  - ▶ Include directories: Add SDL's include folder
  - ▶ Library directories: Add SDL's library folder
  - ▶ Libraries: Add SDL2.lib and windowscodecs.lib
  - ▶ Add the framework code to your project
- ▶ When you run it, you should get a window with a GL rendering + FPS counter

# Program

- ▶ Right now, the program just draws an image to the screen using traditional rasterization techniques
- ▶ We'll use a compute shader to alter the rendered image
- ▶ But first, we must see what CS's look like

# Compute Shader

- ▶ Basic structure of CS:

```
layout(local_size_x=1,local_size_y=1,local_size_z=1) in;  
void main(){  
    uvec3 mynum = gl_GlobalInvocationID;  
    ...do something...  
}
```

## Invocations

- ▶ Since many graphics problems naturally have 2D or 3D structure, API includes built-in support for this
  - ▶ Allows us to organize our code to reflect logical structure of problem
- ▶ Each compute shader invocation has a “position” in a virtual “3D space”
- ▶ When we run CS, we must tell GPU how many invocations to do along each axis of this 3D space

## Dispatch

- ▶ Example: Suppose we've created a Program object with our compute shader code
- ▶ To use it:  
`prog->use();`  
`prog->dispatch(100,1,1);`
  - ▶ This just calls `glDispatchCompute(100,1,1);`
- ▶ Will run compute shader 100 times
  - ▶ `mynum` (in CS) will have values `(0,0,0), (1,0,0), (2,0,0), ... , (99,0,0)`

# Dispatch

- ▶ What if we do:  
    `prog->use();`  
    `prog->dispatch(10,10,1);`
- ▶ This will also run the compute shader 100 times
  - ▶ mynum will have values (0...9,0...9,0)

## Question

- ▶ What's the `local_size` text at the top of the compute shader?



# Workgroup

- ▶ GPU doesn't run one shader invocation at a time
  - ▶ Not efficient
  - ▶ Also, might want several invocations to communicate with each other
- ▶ **Concept: Workgroup**
  - ▶ Collection of several CS invocations
- ▶ Parameters to dispatch are really the number of *workgroups* to be executed
- ▶ Workgroups are defined in a 3D space as well

## Example

- ▶ Suppose workgroup size is (8,1,1)
- ▶ And we call `dispatch(4,1,1)`
- ▶ Total number of CS invocations: 32
  - ▶ `mynum` gets values (0,0,0)...(31,0,0)

## Example

- ▶ Or: If workgroup size is (2,4,8) and we dispatch(3,3,3)
- ▶ Total of  $(3*2)*(3*4)*(3*8) = 1,728$  invocations
  - ▶ mynum.x ranges from 0...5
  - ▶ mynum.y ranges from 0...11
  - ▶ mynum.z ranges from 0...23

## Note

- ▶ Parameters passed to `dispatch()` must be  $\leq 65535$
- ▶ Workgroup size:  $x, y$  must be  $\leq 1024$ ;  $z$  must be  $\leq 64$ 
  - ▶ And  $WG\ x \times WG\ y \times WG\ z$  must be  $\leq 1024$
- ▶ If problem is larger: Call `dispatch()` several times
- ▶ Workgroups usually get mapped to *warps* or *wavefronts* in implementation-defined way

## Note

- ▶ If workgroup size is 1, entire warp is devoted to one thread execution
- ▶ Result: *Thread occupancy* is very low
- ▶ Let's look at an example now

## Example

- ▶ We'll need to alter the files “Globals.h”, “setup.h”, and “draw.h”
- ▶ The rest of the program can remain unchanged
- ▶ First task: Draw to a texture and then copy the texture to the screen

## Create

- ▶ In Globals.h: Add a variable for the FBO:  
`std::shared_ptr<Framebuffer> fbo;`
- ▶ Remember from 2802: FBO allows us to render to a texture instead of the screen

## Create

- ▶ In `setup()`: Initialize the FBO
- ▶ This should be *after* setup creates the globs object:  
globs->fbo =  
`std::make_shared<Framebuffer>(winwidth,winheight,1,GL_RGBA8);`



## Use

- ▶ To use the FBO: We need to alter the draw function:
  - ▶ Rename the existing draw() function to “draw2()”
  - ▶ Add a new draw function:

```
void draw(){
    globs->fbo->setAsRenderTarget(false);
    draw2();
    globs->fbo->unsetAsRenderTarget();
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT );
    glBindFramebuffer( GL_READ_FRAMEBUFFER, globs->fbo->fbo );
    glBindFramebuffer( GL_DRAW_FRAMEBUFFER, 0 );
    glBlitFramebuffer(0,0,globs->fbo->w,globs->fbo->h,
                      0,0,globs->fbo->w,globs->fbo->h,
                      GL_COLOR_BUFFER_BIT,
                      GL_NEAREST );
}
```

## Run It

- ▶ If you run the program, the output should look the same as before
  - ▶ On my system, it was slower: 859 fps (instead of 1017 fps)

## CS

- ▶ Now we're ready to incorporate our compute shader
- ▶ Add a global to Globals.h:  
`Program cs{"cs.txt"};`

## Draw

- ▶ Change draw to use the CS
- ▶ After the `fbo->unsetRenderTarget()` but before the `glClear`: Add some code:

```
globs->cs.use();  
globs->fbo->texture->bindImage(0);  
globs->cs.dispatch(globs->fbo->w, globs->fbo->h, 1 );  
glMemoryBarrier( GL_ALL_BARRIER_BITS );  
globs->fbo->texture->unbindImage(0);
```

- ▶ Finally, the compute shader itself: in shaders/cs.txt:

```
layout(local_size_x=1,local_size_y=1,local_size_z=1) in;
layout(binding=0,rgba8) uniform image2DArray img2;
void main(){
    ivec3 mynum = ivec3(gl_GlobalInvocationID);
    vec4 c = imageLoad( img2, mynum );
    if( length(c.rgb) < 0.1 ){
        c.rgb = vec3(0,1,0);
        imageStore(img2, mynum, c );
    }
}
```

## Explanation

- ▶ CS loads a single texel from the texture bound to image unit #0
- ▶ If it is close to black, we change the color to green
- ▶ Finally, we store the value

## Results

- ▶ On my machine, I get these results:
  - ▶ Original code: 859 frames per second
  - ▶ CS code: 92 frames per second
- ▶ Wow! Why so low?

## Problem

- ▶ Notice our workgroup size: We are using only one thread of every warp
  - ▶ I ran this on an nVidia card, so warp size is 32
  - ▶ On AMD, wavefront size is 64
- ▶ So only  $1/32 = 3\%$  of GPU being used
- ▶ What if we change the code...



## Changes

- ▶ Change first line of CS:  
`layout(local_size_x=64,local_size_y=1,local_size_z=1) in;`
- ▶ And change `draw()`:  
`globs->cs.dispatch(globs->fbo->w/64, globs->fbo->h, 1 );`

## Results

- ▶ No CS: 859 frames per second
- ▶ CS, workgroup size=1: 92 frames per second
- ▶ CS, workgroup size=64: 727 frames per second
- ▶ That's better...But still only 84% of original results
  - ▶ Is it really *that* costly to change some pixel colors?

## Reasons

- ▶ Several things combine to make CS version of program slower
- ▶ How many operations can GPU do per second?
  - ▶ Switch render target (60,000 per sec)
  - ▶ Switch program (300,000/sec)
  - ▶ Switch pipeline state (glEnable/glDisable) (maybe 700,000/sec)
  - ▶ Switch textures (1,500,000/sec)
  - ▶ Switch VAO/UBO bindings (maybe 5,000,000/sec)
  - ▶ Update uniform (10,000,000/sec)
- ▶ Switching to/from CS can be more costly than just switching VS/FS programs

## Test

- ▶ How can we determine overhead of using CS?
- ▶ Can you think of a way (without looking at following slides?)

## Overhead

- ▶ One option: Comment out the `csprog->dispatch()` line in `draw()`
  - ▶ Result: 870 FPS
  - ▶ This is essentially the same as the no-cs result
    - ▶ It's actually faster!
    - ▶ Probably OS/scheduler overhead
- ▶ I suspect the GPU notices the CS is never dispatched and doesn't do any of the other setup work
  - ▶ Note: Explain concept of GPU ring buffer

## Overhead

- ▶ Next test: Put `csprog->dispatch()` back in and alter the shader itself
- ▶ Make first line be “return”
- ▶ Result: 835 FPS
- ▶ Analysis: What does this tell you?

## Analysis

- ▶ This tells us the cost of using the CS
- ▶ Again, the numbers:
  - ▶ No dispatch: 859 FPS
  - ▶ Dispatch, but CS returns immediately: 835 FPS
  - ▶ CS does real work: 727 FPS
- ▶ We pay some cost just to have a CS
- ▶ Most of the hit comes from the work the CS is doing

# The Birds and the Bees

- ▶ Or, where little pixels come from...



# SIMD

- ▶ Remember: We said GPU was SIMD
- ▶ Modern chipset manufacturers take this to the extreme
- ▶ Consider some example shader code...

# Shader

- ▶ Typical Lambertian shader:

```
layout(binding=0) uniform sampler2D tex;
uniform vec3 lightPosition;
in vec2 v_texCoord;
in vec3 v_normal;
in vec3 v_worldPos;
out vec4 color;
void main(){
    vec4 c = texture(tex, v_texCoord);
    vec3 L = normalize(lightPosition - v_worldPos);
    vec3 N = normalize(v_normal);
    float dp = dot(L,N);
    dp = max(0.0, dp);
    color.rgb = dp * c.rgb;
    color.a = 1.0;
}
```

- ▶ What does this look like at GPU instruction level?

# Code

- ▶  

```
;format: INSTRUCTION dest, args
SUB tmp1.xyz, uniform[0].xyz, input[2].xyz ;lightPos - worldPos
DOT tmp2.x, tmp1.xyz, tmp1.xyz             ;length squared
RSQRT tmp2.x, tmp2.x                       ;one over length
MUL tmp1.xyz, tmp2.xxx, tmp1.xyz           ;tmp1 gets unit L
DOT tmp2.x, input[1].xyz, input[1].xyz     ;get sq. len of normal
RSQRT tmp2.x, tmp2.x                       ;one over length
MUL tmp2.xyz, tmp2.xxx, input[1].xyz       ;tmp2 holds unit N
DOT tmp1.x, tmp1.xyz, tmp2.xyz             ;tmp1 gets dot(L,N)
MAX tmp1.x, tmp1.x, 0                      ;tmp1 holds max(dot,0)
TEXFETCH tmp0, 0, input[0]                ;fetch from tex unit 0
MUL output.xyz, tmp1.xxx, tmp0.xyz         ;dp * c.rgb
MOV output.w, 1.0
```
- ▶ Notice: Compiler moved texture fetch as far down as it could.  
Why?

## GPU

- ▶ GPU starts executing shader code
- ▶ All threads in warp/wavefront proceed in tandem
  - ▶ Same instructions, but different data
- ▶ They're all busy...until they hit TEXFETCH
- ▶ This requires going out to memory
  - ▶ Up to this point, all data was in local registers

# Memory

- ▶ Memory access is slow
  - ▶ Even though it's local, dedicated RAM on GPU
  - ▶ Still requires arbitrating GPU bus, putting request out, waiting for response to come back
- ▶ Solution: GPU:
  - ▶ Queues up the TEXFETCH requests (one per thread in warp/wavefront)
  - ▶ Puts all 32 (64) threads on ice (suspends them)
  - ▶ Starts executing FS on brand new batch of 32 (64) pixels
- ▶ When those FS's hit TEXFETCH, put them on ice and repeat above process

## Result

- ▶ Eventually, GPU will run out of resources for suspending threads
  - ▶ The freezer gets full!
- ▶ So it pulls a batch out of suspended animation
  - ▶ Hopefully, results of TEXTFETCH are available by now
- ▶ Begin executing
- ▶ When batch is finished, grab a new batch of pixels (or un-suspend a waiting group)

## Code

- ▶ This explains why compiler moved TEXTFETCH down
- ▶ Want to make maximum possibility of overlapping work
- ▶ Imagine TEXTFETCH was at top of code
  - ▶ We'd queue up the fetch, switch threads, queue another fetch, etc.
  - ▶ No overlap of memory fetch + computation
- ▶ Moral: CS's are best used if we have computation-heavy work to do

## CS

- ▶ Now we can return to our compute shader
  - ▶ Wasn't doing much computation
  - ▶ So most of the time, GPU is just queueing up memory read/write operations
- ▶ Why was this not a problem with FS?
  - ▶ CS's are more flexible than FS's
  - ▶ So GPU can make assumptions regarding what will and won't happen when it's doing FS's



## Detail

- ▶ One more detail: What's the `glMemoryBarrier()`?
  - ▶ Tells GPU that we need to make sure that all previous memory writes have completed before we go on to following code
- ▶ CS writes are *incoherent* by default
  - ▶ Not visible/published to rest of system until barrier is hit
    - ▶ Except for some frankly hard-to-remember exceptions...
- ▶ Best to put it in to make sure you get correct results

## Assignment

- ▶ Use the compute shader to compute a quarter size version of the original rendered scene.
- ▶ Draw the quarter-size version to the lower left quadrant of the window.
- ▶ The rest of the window should be black.

# Sources

- ▶ Khronos Corp. OpenGL reference pages & quick reference card.
- ▶ [https://en.wikipedia.org/wiki/List\\_of\\_Nvidia\\_graphics\\_processing\\_units](https://en.wikipedia.org/wiki/List_of_Nvidia_graphics_processing_units)
- ▶ [https://asteroidsathome.net/boinc/cpu\\_list.php](https://asteroidsathome.net/boinc/cpu_list.php)
- ▶ [https://www.khronos.org/opengl/wiki/Compute\\_Shader](https://www.khronos.org/opengl/wiki/Compute_Shader)
- ▶ Cass Everitt and John McDonald. Beyond Porting.  
<http://media.steampowered.com/apps/steamdevdays/slides/beyondporting.pdf> (talk given in 2014)
- ▶ <https://www.slideshare.net/CassEveritt/approaching-zero-driver-overhead>
- ▶ [https://www.khronos.org/opengl/wiki/Shader\\_Storage\\_Buffer\\_Object](https://www.khronos.org/opengl/wiki/Shader_Storage_Buffer_Object)
- ▶ [https://www.khronos.org/opengl/wiki/Memory\\_Model#Incoherent\\_memory\\_access](https://www.khronos.org/opengl/wiki/Memory_Model#Incoherent_memory_access)
- ▶ <https://www.khronos.org/opengl/wiki/GLAPI/glDeleteSync>
- ▶ [https://www.khronos.org/opengl/wiki/Sync\\_Object](https://www.khronos.org/opengl/wiki/Sync_Object)
- ▶ [https://www.khronos.org/opengl/wiki/Synchronization#Implicit\\_synchronization](https://www.khronos.org/opengl/wiki/Synchronization#Implicit_synchronization)
- ▶ <https://www.khronos.org/opengl/wiki/GLAPI/glMapBufferRange>
- ▶ <https://www.khronos.org/opengl/wiki/GLAPI/glBufferStorage>
- ▶ [https://www.khronos.org/opengl/wiki/Buffer\\_Object\\_Streaming](https://www.khronos.org/opengl/wiki/Buffer_Object_Streaming)
- ▶ [https://www.khronos.org/opengl/wiki/Buffer\\_Object](https://www.khronos.org/opengl/wiki/Buffer_Object)
- ▶ <http://www.equasys.de/colorconversion.html>
- ▶ [https://docs.opencv.org/3.1.0/de/d25/imgproc\\_color\\_conversions.html](https://docs.opencv.org/3.1.0/de/d25/imgproc_color_conversions.html)
- ▶ John D. Cook. Converting color to grayscale.  
<https://www.johndcook.com/blog/2009/08/24/algorithms-convert-color-grayscale/>

Created using L<sup>A</sup>T<sub>E</sub>X.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>