

Compute Shaders 2

Review

- ▶ Purpose
- ▶ Coding
 - ▶ C++ / GLSL

Motivation

- ▶ We've seen one way to get data to/from GPU: Textures (Images)
 - ▶ Well suited for image-based operations
 - ▶ Easy to use and integrate with existing render pipeline
- ▶ But some important limitations
 - ▶ Only a limited number of image units available
 - ▶ Often, it's limited to 8
 - ▶ Only suited for homogeneous data
 - ▶ Limited selection: float, vec{2,3,4}
 - ▶ Not easy to read back data to CPU

Buffers

- ▶ We can solve some issues with buffers
- ▶ Usually, buffers used for specifying vertex/index data
- ▶ But they are more flexible than that!

Buffers

- ▶ Recall: GPU has several buffer binding points
 - ▶ Attach buffer to binding point, work with it
 - ▶ Some binding points are indexed: Can attach different buffers to different indices
- ▶ Ex: For drawing
 - ▶ Attach buffer with vertex data to `GL_ARRAY_BUFFER` binding point
 - ▶ Attach buffer with vertex indices to `GL_ELEMENT_ARRAY_BUFFER` binding point
- ▶ For compute shaders, most interesting binding point is `GL_SHADER_STORAGE_BUFFER`

Usage

- ▶ We have framework for buffers in Buffer.h file
- ▶ If we want to communicate data CPU \leftrightarrow GPU: Create *mappable* buffer:
- ▶ But there are (as usual) some complications...

Code

- ▶ Compute shader that does (very simple) buffer operation:

```
layout(local_size_x=64,local_size_y=1,local_size_z=1) in;  
layout(std430,binding=0,row_major) buffer Foo {  
    float buff[];  
};  
void main(){  
    ivec3 mynum = ivec3(gl_GlobalInvocationID);  
    buff[mynum[0]] = mynum[0]*10;  
}
```

- ▶ Note buffer declaration syntax

- ▶ std430 = Packing rules
- ▶ binding = Which binding point
- ▶ row_major = For matrices

Code

- ▶ Now for the CPU code: [main.cpp](#)
- ▶ When I run it, I get this output:
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 - ▶ That doesn't look right!

Problem

- ▶ GPU hasn't finished when CPU tries to access memory
- ▶ We need to synchronize the two sides
- ▶ Better code: [main.cpp](#)
- ▶ Now I get this output:
0 10 20 30 40 50 60 70 80 90 100 110 120 130 140 150
- ▶ Much better!

Explanation

- ▶ `glMemoryBarrier(GL_ALL_BARRIER_BITS)`
 - ▶ This tells GL to insert a *memory barrier*: Things written before the barrier will be visible on the CPU when we sync on the data
- ▶ `auto sync =`
`glFenceSync(GL_SYNC_GPU_COMMANDS_COMPLETE,0);`
 - ▶ This creates a *fence* object
 - ▶ Fences allow us to partition commands into two groups: One for each side of the fence
- ▶ `auto rv = glClientWaitSync(sync,`
`GL_SYNC_FLUSH_COMMANDS_BIT, -1);`
 - ▶ This tells CPU to wait if the fence has not been reached by GPU
 - ▶ Third argument = timeout (nanoseconds, 64 bit unsigned value)
- ▶ `glDeleteSync(sync);`
 - ▶ Frees resources used by the fence

Note

- ▶ We don't need `glMemoryBarrier` and fence sync unless we plan to read back data *from GPU to CPU*
- ▶ If all data stays on GPU, no need for these

Structured Buffers

- ▶ Inconvenient to have only primitive types in buffers
- ▶ Example: Suppose we are processing collection of spheres
- ▶ Each sphere has:
 - ▶ Center (vec3)
 - ▶ Radius (float)
 - ▶ Color (vec4)
 - ▶ Reflectivity (float)

SoA

- ▶ We could code using "Structure of Arrays" (SoA) format:

```
layout(binding=0) buffer Foo {  
    vec3 center[];  
};  
layout(binding=1) buffer Foo2 {  
    float radius[];  
};  
layout(binding=2) buffer Foo3 {  
    vec4 color[];  
};  
layout(binding=3) buffer Foo4 {  
    float reflectivity[];  
};
```

SoA

► To process a sphere:

```
uint idx = ...;  
vec3 c = center[idx];  
float r = radius[idx];  
vec4 co = color[idx];  
float ref = reflectivity[idx];  
...use variables...
```

Problem

- ▶ One problem: Data for one sphere is “spread out” over memory
 - ▶ When we load data (CPU side): Must distribute it to different buffers
 - ▶ This may require extra time for conversion / data copying
- ▶ Also: We're limited in how many buffers we can have active at once
 - ▶ Many GPU's only permit 8
- ▶ We can use AoS format to get around this

Structs

- ▶ Defining structs: We need to do this in GLSL and in C
- ▶ Ex:

```
//GLSL
struct Sphere {
    vec4 center;
    float radius;
    vec4 color;
    uint reflectivity;
};
layout(std430,binding=0) buffer Foo {
    Sphere spheres[];
}
void main(){
    ...
    vec3 c = spheres[idx].center.xyz;
}
```


Padding

- ▶ GPU keeps data padded
- ▶ Alignment: A variable of a given type has an alignment (often, but not always, equal to its size)
 - ▶ Floats, ints, uints = 4
 - ▶ Doubles = 8
 - ▶ vec2, ivec2, uvec2 = 8
 - ▶ dvec2 = 16
 - ▶ vec3, ivec3, uvec3, vec4, ivec4, uvec4 = 16
 - ▶ dvec3, dvec4 = 32
- ▶ Arrays of items have each item one after another
- ▶ An MxN matrix is treated as an M element array of vecN's

Padding

- ▶ If an object would land at an impermissible address, it gets pushed further along until its address is OK
- ▶ In the case of the entire struct, it will have trailing padding so the size of the struct is a multiple of the largest n of any element in the struct
- ▶ Some examples may help...

Example

▶ GLSL declaration:

```
struct Foo {  
    int a;  
    int b;  
};
```

▶ Memory layout:

- ▶ 4 bytes for a
 - ▶ 4 bytes for b
- ▶ Total size of structure: 8 bytes
 - ▶ No trailing padding

▶ C++ declaration:

```
struct Foo {  
    alignas(4) int32_t a;  
    alignas(4) int32_t b;  
};
```

Example

- ▶ GLSL declaration:

```
struct Foo {  
    float a;  
    double b;  
};
```

- ▶ Memory layout:

- ▶ 4 bytes for a
 - ▶ 4 bytes of padding
 - ▶ 8 bytes for b
- ▶ Total size of structure: 16 bytes
 - ▶ No trailing padding

- ▶ C++ declaration:

```
struct Foo {  
    alignas(4) float a;  
    alignas(8) double b;  
};
```

Example

► GLSL declaration:

```
struct Foo {  
    double a;  
    uint b;  
};
```

► Memory layout:

- 8 bytes for a
 - 4 bytes for b
 - Total size: 12 bytes
 - Largest n: 8
 - So 4 bytes of padding
- Total size of structure: 16 bytes

► C++ declaration:

```
struct Foo {  
    alignas(8) double a;  
    alignas(4) uint32_t b;  
};
```

Example

► GLSL declaration:

```
struct Foo {  
    int a[3];  
    int b[4];  
};
```

► Memory layout:

- 12 bytes for a (4 per int)
 - 16 bytes for b (4 per int)
- Total size of structure: 28 bytes
 - No padding

► C++ declaration:

```
struct Foo {  
    alignas(4) int32_t a[3];  
    alignas(4) int32_t b[4];  
};
```

Example

- ▶ GLSL declaration:

```
struct Foo {  
    double a[3];  
    int b[4];  
};
```

- ▶ Memory layout:

- ▶ 24 bytes for a (8 per double)
 - ▶ 16 bytes for b (4 per int)
- ▶ Total size of structure: 40 bytes
- ▶ No padding needed

- ▶ C++ declaration:

```
struct Foo {  
    alignas(8) double a[3];  
    alignas(4) int32_t b[4];  
};
```

Example

- ▶ GLSL declaration:

```
struct Foo {  
    double a[3];  
    int b;  
};
```

- ▶ Memory layout:

- ▶ 24 bytes for a (8 per double)
 - ▶ 4 bytes for b
 - ▶ Total size = 28. Largest n = 8
 - ▶ So 4 bytes for padding
- ▶ Total size of structure: 32 bytes

- ▶ C++ declaration:

```
struct Foo {  
    alignas(8) double a[3];  
    alignas(4) int32_t b;  
};
```


Example

▶ GLSL declaration:

```
struct Foo {  
    vec4 a;  
    int b;  
};
```

▶ C++ declaration:

```
struct Foo {  
    alignas(16) vec4 a;  
    alignas(4) int32_t b;  
};
```

▶ Memory layout:

- ▶ 16 bytes for a (4 per element)
 - ▶ 4 bytes for b
 - ▶ Total = 20. Largest alignment $n = 16$ (for the vec4).
 - ▶ 12 bytes of padding to make it 32
- ▶ Total size of structure: 32 bytes

Example

► GLSL declaration:

```
struct Foo {  
    dvec4 a;  
    int b;  
};
```

► Memory layout:

- 32 bytes for a (8 per element)
 - 4 bytes for b
 - Total size: 36
 - Largest $n = 4 * 8 = 32$ (from the dvec4)
 - Next highest multiple of 32 is 64
 - Thus, 28 bytes of padding ($36 + 28 = 64$)
- Total size of structure: 64 bytes

► C++ declaration:

```
struct Foo {  
    //math3d doesn't have dvec  
    alignas(8) double a[4];  
    alignas(4) int32_t b;  
};
```

Example

- ▶ GLSL declaration:

```
struct Foo {  
    int a;  
    dvec4 b;  
};
```

- ▶ Memory layout:

- ▶ 4 bytes for a
- ▶ 28 bytes of padding
- ▶ 32 bytes for b
- ▶ Total size: 64. No trailing padding.

- ▶ Total size of structure: 64 bytes

- ▶ C++ declaration:

```
struct Foo {  
    alignas(4) int32_t a;  
    alignas(8) double b[4];  
};
```

Example

► GLSL declaration:

```
struct Foo {  
    int a;  
    dvec4 b;  
    int c;  
};
```

► C++ declaration:

```
struct Foo {  
    alignas(4) int32_t a;  
    alignas(32) double b[4];  
    alignas(4) int32_t b;  
};
```

► Memory layout:

- 4 bytes for a, 28 bytes of padding, then 32 bytes for b, 4 bytes for c
- Total size: $4+28+32+4 = 68$
- Largest $n = 32$ (for the dvec4). Next highest multiple of 32 is 96
- So $(96-68) = 28$ bytes of trailing padding
- Total size of structure: 96 bytes

Example

► GLSL declaration:

```
struct Foo {  
    vec3 a;  
    mat3x2 b;  
};
```

► Memory layout:

- 12 bytes for vec3
- 4 bytes of padding
- b is treated as if it were declared as vec2 b[3]
- vec2's must start on an 8 byte boundary
- Then 24 bytes for b (8 bytes per vec3, 3 of them)
- Total: 40 bytes. Largest n = 16 (for the vec3), so must have 8 more bytes of padding to make 48.

► C++ declaration:

```
struct Foo {  
    alignas(16) vec3 a;  
    alignas(8) vec2 b[3];  
};
```

Example

▶ GLSL declaration:

```
struct Foo {  
    vec3 a;  
    mat3 b;  
};
```

▶ C++ declaration:

```
struct Foo {  
    alignas(16) vec3 a;  
    alignas(16) vec4 b[3];    //  
    Notice vec4!  
};
```

▶ Memory layout:

- ▶ 12 bytes for a, 4 bytes of padding, 48 bytes for b
 - ▶ Each vec3 takes up $3 \times 4 = 12$ bytes
 - ▶ 4 bytes of padding after each so next vec3 starts on address divisible by 16
- ▶ Total: 64 bytes. No trailing padding needed.

Example

- ▶ Going back to the sphere...

```
//GLSL
struct Sphere {
    vec4 center;
    float radius;
    vec4 color;
    float reflectivity;
};
```

► Going back to the sphere:

```
struct Sphere {  
    alignas(16) vec4 center;  
    alignas(4) float radius;  
    alignas(16) vec4 color;  
    alignas(4) float reflectivity;  
};
```


One More Example

- ▶ Suppose we have some sort of particles with position and velocity
- ▶ GLSL:

```
struct ParticleData {  
    vec3 position;  
    vec3 velocity;  
};
```

- ▶ C:

```
struct ParticleData {  
    alignas(16) vec3 position;  
    alignas(16) vec3 velocity;  
};
```

Example

- ▶ Now we can see an example of using CS to do some more intensive work
- ▶ Example: Sphere-based raytracer

Code

- ▶ Start with the GL framework code
- ▶ Replace some files:
 - ▶ [draw.h](#)
 - ▶ [Globals.h](#)
 - ▶ [setup.h](#)
 - ▶ [shaders/cs.txt](#)
- ▶ If you run this, you should get a solid magenta screen

Code

- ▶ We have several things we need to do:
 - ▶ Get the spheres to the compute shader
 - ▶ Do the raytracing in the compute shader

Data

- ▶ First we need to get the data to a buffer so the CS can use it
- ▶ Define a type at the top of setup.h:

```
struct GPUSphere{  
    alignas(16) vec4 centerAndRadius;  
    alignas(16) vec4 color;  
};
```

Data

- ▶ Add some code to the bottom of setup():

```
std::vector<GPUSphere> sphereData(globs->scene.spheres.size());
for(unsigned i=0;i<globs->scene.spheres.size();++i ){
    sphereData[i].centerAndRadius = vec4(
        globs->scene.spheres[i].c, globs->scene.spheres[i].r);
    sphereData[i].color = vec4( globs->scene.spheres[i].color, 1.0
        );
}
auto b = Buffer::create(sphereData);
b->bindBase(GL_SHADER_STORAGE_BUFFER, 0);
```

Data

- ▶ Also set some uniforms in `setup()`:

```
Program::setUniform("lightPosition", globs->scene.lightPosition);  
Program::setUniform("lightColor", vec3(1,1,1) );
```

CS

- ▶ In the CS, we define the Sphere type and the buffer that holds the data:

```
struct Sphere {  
    vec4 centerAndRadius;  
    vec4 color;  
};  
layout(std430, binding=0) buffer S {  
    Sphere spheres[];  
};
```


Draw

- ▶ We change the draw.h file: Dispatch work to CS, then blit result to screen
- ▶ Code: [draw.h](#)

- In CS, we do exactly one ray per invocation. Outline:

```
void main(){
    ivec2 pixelCoord = ivec2(gl_GlobalInvocationID.xy);
    ivec2 picSize = imageSize(img).xy;
    float d = 1.0 / 0.69974612;      //denom = tan(35 degrees) = field of view
    float dy = 2.0/(picSize.y-1);
    float dx = 2.0/(picSize.x-1);
    float y = -1.0 + pixelCoord.y * dy;
    float x = -1.0 + pixelCoord.x * dx;
    vec3 rayDir = x*cameraRight + y*cameraUp + d*cameraLook;
    bool wasIntersection;
    vec3 ip, N, color;
    wasIntersection = traceSpheres(eyePos, rayDir, ip, N, color );
    if( !wasIntersection )
        color = vec3(0,0,0);
    else
        color = shadePixel( ip, N, color );
    imageStore(img, ivec3(pixelCoord,0), vec4(color,1.0) );
}
```

That's It!

- ▶ That's it!
 - ▶ (Oh, except for changing Globals to load [scene3.txt](#) instead)

Results

- ▶ For spheres (scene3.txt):
 - ▶ CPU raytracing (Core i7, 2.4GHz): 0.978 sec/frame
 - ▶ GPU raytracing (Intel HD 5500): 0.084 sec/frame
 - ▶ GPU raytracing (GeForce 920M): 0.043 sec/frame

Assignment

- ▶ Implement triangle-mesh raytracing on the GPU using a compute shader
- ▶ For reference, here's what I got for scene2.txt:
 - ▶ CPU: 8.60 sec/frame
 - ▶ GPU (Intel HD 5500): 0.575 sec/frame
 - ▶ GPU (GeForce 920M): 0.290 sec/frame

Sources

- ▶ [https://www.khronos.org/opengl/wiki/ Shader_Storage_Buffer_Object](https://www.khronos.org/opengl/wiki/Shader_Storage_Buffer_Object)
- ▶ [https://www.khronos.org/opengl/wiki/ Memory_Model#Incoherent_memory_access](https://www.khronos.org/opengl/wiki/Memory_Model#Incoherent_memory_access)
- ▶ [https://www.khronos.org/opengl/wiki/ GLAPI/glDeleteSync](https://www.khronos.org/opengl/wiki/GLAPI/glDeleteSync)
- ▶ [https://www.khronos.org/opengl/wiki/ Sync_Object](https://www.khronos.org/opengl/wiki/Sync_Object)
- ▶ [https://www.khronos.org/opengl/wiki/ Synchronization#Implicit_synchronization](https://www.khronos.org/opengl/wiki/Synchronization#Implicit_synchronization)
- ▶ [https://www.khronos.org/opengl/wiki/ GLAPI/glMapBufferRange](https://www.khronos.org/opengl/wiki/GLAPI/glMapBufferRange)
- ▶ [https://www.khronos.org/opengl/wiki/ GLAPI/glBufferStorage](https://www.khronos.org/opengl/wiki/GLAPI/glBufferStorage)
- ▶ [https://www.khronos.org/opengl/wiki/ Buffer_Object_Streaming](https://www.khronos.org/opengl/wiki/Buffer_Object_Streaming)
- ▶ [https://www.khronos.org/opengl/wiki/ Buffer_Object](https://www.khronos.org/opengl/wiki/Buffer_Object)

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>