

Compute Shaders 4

Motivation

- ▶ Go faster!

Recall

- ▶ Recall our benchmarks for raytracing triangle mesh:
 - ▶ CPU (basic algorithm): 8.60 sec/frame
 - ▶ CPU (octrees): 0.090 sec/frame
 - ▶ GPU (basic algorithm): 0.043 sec/frame
- ▶ CPU + acceleration structure is almost as good as GPU!
- ▶ Can we use an acceleration structure on the GPU?

Acceleration

- ▶ We designed our octree traversal algorithm to be nonrecursive, so that's not an issue
- ▶ But: Recall the definition of an octree node:

```
class OctreeNode{
public:
    vec3 min, max;          //bounding volume; only used when
                           building
    std::array<unsigned,8> children;    //0=no child
    std::vector<Triangle> triangles;
    std::array<vec4,6> planes;
    static std::vector<OctreeNode> nodes;
    ...
};
```

Problem

- ▶ The problem we'll face is that an octree node has a variable-length data structure: The triangles list
 - ▶ All the rest are fixed-size and thus easy to convert

Recall

- ▶ What data do triangles have?

```
class Triangle{  
    public:  
        vec3 p[3];           //vertices  
        vec3 N;              //normal, unit length  
        float D;             //plane equation D  
};
```

- ▶ GLSL:

```
struct Triangle {  
    vec4 p[3];  
    vec4 ND;           //xyz=normal, w=D  
};  
layout(std430, binding=0) buffer Oct {  
    octreeNode[];  
};
```

Octree Node

► Octree node: In GLSL:

```
struct OctreeNode {  
    uint children[8];           //indices in octreeNodes  
    ??? triangles;              //what to do here?  
    vec4 planes[6];             //ABCD of each plane  
};  
//notice: binding = 1!!!  
layout(std430,binding=1) buffer Oct {  
    octreeNodes[];  
};
```

Pattern

- ▶ We can use a pair of integers to hold “pointers” to a list of triangles
 - ▶ First pointer = index in the list of first triangle
 - ▶ Second = index past last triangle
- ▶ So we have:

```
struct OctreeNode {  
    uint children[8];           //indices in octreeNode  
    uint firstTriangle;        //index in triangles[]  
    uint lastTriangle;  
    vec4 planes[6];           //ABCD of each plane  
};
```


Note

- ▶ Note that CPU side code needs to rearrange triangles so they are contiguous
 - ▶ If we just use the Mesh data as it comes in: Won't work!
- ▶ Example updated code: [Octree.h](#)

Compute Shader

- Now we can incorporate the octree code in the compute shader:

```
bool traceOctree(vec3 s, vec3 v, out vec3 ip, out vec3 N, out vec3 color ){
    float closestT = -1.0;
    bool inter=false;
    uint stk[128];
    uint top=0;
    stk[top++] = 0; //push 0
    while( top != 0 ){
        uint ni = stk[--top];    //pop
        OctreeNode node = octreeNodes[ni];
        if( rayBoxIntersection( node.planes, s, v ) ){
            if( node.firstTriangle == node.lastTriangle ){
                for(int i=0;i<8;++i){
                    uint ci = node.children[i];
                    if( ci != 0 )
                        stk[top++] = ci;    //push(ci)
                }
            } else {
                //leaf
                if( traceTriangles( node.firstTriangle, node.lastTriangle, s,v,ip,N,closestT,color) )
                    inter=true;
            }
        }
    }
    return inter;
}
```

traceTriangles

- ▶ Almost identical to what we had before:

```
bool traceTriangles(uint first, uint last, vec3 s, vec3 v, out
    vec3 ip, out vec3 N,
    inout float closestT, out vec3 color )
{
    bool inter=false;
    for(uint i=first;i<last;++i){
        Triangle T = triangles[i];
        ...do intersection tests...
    }
}
```

Results

- ▶ CPU (basic algorithm): 8.60 sec/frame
- ▶ CPU (octrees): 0.090 sec/frame
- ▶ GPU (basic algorithm): 0.043 sec/frame
- ▶ GPU (Intel HD 5000): 0.011 sec/frame
- ▶ GPU (GeForce 920M): 0.024 sec/frame

Memory Organization

- ▶ Up to now, we haven't paid much attention to GPU storage types
- ▶ But it actually can make a big difference!
- ▶ What kind of storage is there?

Textures

- ▶ Textures
 - ▶ Potentially large amount of storage: Usually at least 4096×4096 (=16M texels)
 - ▶ Read-only (if attached to sampler) or Read/write (if attached to image unit)
 - ▶ Number available: Fairly high for samplers (Modern hardware = 96); fairly low for image units (8)
 - ▶ Homogeneous type; scalar or $\text{vec}\{2,3,4\}$
 - ▶ Comparatively slow to access (although GPU caching helps sequential access)

Buffers

- ▶ Buffers
 - ▶ Larger amount of storage: Often, up to entire size of GPU RAM (ex: Gigabyte range)
 - ▶ Read/write
 - ▶ Limited number available (8)
 - ▶ Can use structures to define heterogeneous types
 - ▶ Comparatively slow to access (although GPU caching helps sequential access)

Uniforms

- ▶ Uniforms
 - ▶ Limited space: Maybe as small as 16KB
 - ▶ Read only
 - ▶ Can use structures to define heterogeneous types
 - ▶ Fast access

Variables


- ▶ Global/local variables
 - ▶ Number available: Depends on specific hardware
 - ▶ Read/write
 - ▶ Private to one shader invocation
 - ▶ Fast access

Problem

- ▶ We've been using buffers for our data
 - ▶ We probably have too much for uniforms
- ▶ But: Buffers aren't the fastest way to access memory

Architecture

- Recall GPU architecture:



Texture/Buffer

- ▶ Textures and buffers sit in global RAM
- ▶ All the PE's contend for access to it
- ▶ And: It's not designed to be exceptionally fast (latency)
- ▶ This can be a problem for us...

Code

- ▶ Ignoring octrees, let's go back to the original GPU raytracing code
- ▶ What does it look like?


```
bool traceTriangles(vec3 s, vec3 v, out vec3 ip, out vec3 N,  
                    out vec3 color ){  
    bool inter=false;  
    float closestT = -1.0;  
    int numTris = triangles.length();  
    for(int i=0;i<numTris;++i){  
        Triangle T = triangles[i];  
        float denom = dot(T.ND.xyz,v);  
        ...etc...  
    }  
}
```

Observe

- ▶ The access to `triangles[]` is not very fast (compared to computations)
 - ▶ Remember, it's stored in a buffer object
- ▶ And *all* the thread invocations are contending for access to this buffer
- ▶ It seems like we could improve this somehow...
 - ▶ ...We can!

Architecture

- ▶ A closer look at GPU architecture:



Memory

- ▶ Each workgroup has a local private memory area
- ▶ We can take advantage of this by sharing buffer data between workers

Approach

- ▶ First, we declare a global variable that's shared between the workgroups:
shared Triangle cachedTriangles[64];
- ▶ Next, we adjust traceTriangles:
 - ▶ Each worker then copies exactly one thing from the triangles buffer to the cache
 - ▶ Then all the workers do ray intersections with items in the cache
 - ▶ Since we have more triangles than workers, we repeat the above steps until all triangles have been processed

ID

- ▶ `gl_LocalInvocationID` will give us our workgroup ID number
- ▶ Workgroup size was defined at the very top of the CS. Suppose we had something like this:

```
#define WORKGROUP_SIZE 64  
layout(local_size_x=WORKGROUP_SIZE,local_size_y=1,local_size_z=1)  
in;
```

- ▶ Note: Before, we had size hardcoded at something like 64

Code

► We get something like this:

```
bool traceTriangles(vec3 s, vec3 v, out vec3 ip, out vec3 N, out vec3 color ){
    uint myId = gl_LocalInvocationID.x;
    bool inter=false;
    float closestT = -1.0;
    int numTris = triangles.length();
    for(int k=0;k<numTris;k+=WORKGROUP_SIZE){
        int last = k+WORKGROUP_SIZE;
        if( last > numTris )
            last = numTris;
        int numToCopy = last-k;
        if( myId < numToCopy ){
            cachedTriangles[myId] = triangles[k+myId];
        }
        for(int i=0;i<numToCopy;++i){
            Triangle T = cachedTriangles[i];
            float denom = dot(T.ND.xyz,v);
            ...use T...
        }
    }
}
```

Results

- ▶ Intel HD 5500:
 - ▶ No shared memory cache: 0.576 sec/frame
 - ▶ Shared memory cache: 0.475 sec/frame (1.2x speedup)
- ▶ GeForce 920M:
 - ▶ No shared memory cache: 0.291 sec/frame
 - ▶ Shared memory cache: 0.365 sec/frame (Surprise! It slowed down!)
- ▶ But there's a problem...

Problem

- ▶ The display is corrupted. Why?



Problem

- ▶ We don't know that all threads of workgroup proceed in lockstep
 - ▶ It's "mostly" simultaneous
 - ▶ But some workers can get held up (ex: Data fetch from buffer memory is delayed)
- ▶ Other workers continue on, assuming all data has been put in shared buffer

Solution

- ▶ We need to insert *barrier* to tell all threads to wait until everyone has arrived at a certain point

```
for(int k=0;k<numTris;k+=WORKGROUP_SIZE){  
    int last = k+WORKGROUP_SIZE;  
    if( last > numTris )  
        last = numTris;  
    int numToCopy = last-k;  
  
    barrier();  
    if( myId < numToCopy ){  
        cachedTriangles[myId] = triangles[k+myId];  
    }  
    groupMemoryBarrier();  
    barrier();  
}
```

Explanation

- ▶ `barrier()` halts execution until all threads have reached that point
- ▶ `groupMemoryBarrier` ensures any writes prior to that point are visible to all other threads
- ▶ It makes sense why we'd need these two after the updating of `cachedTriangles`
- ▶ But why do we need `barrier()` *before* we write to `cachedTriangles`?

Explanation

- ▶ Remember, this is in a loop (the for-k loop)
- ▶ Suppose thread A completes the loop and cycles around for the next k-loop iteration
- ▶ If other threads are still working, they will assume that cachedTriangles isn't going to change out from under them
- ▶ But if A could just proceed on, that's exactly what would take place!
 - ▶ Result: We still get screen corruption
- ▶ Thus, we need barrier before and after.
 - ▶ If there was no for-k loop then we'd only need the post-barrier, not the pre-barrier
 - ▶ In my tests, adding the barrier gave correct display and only slowed rendering down by a few msec per frame

Sources

- ▶ https://www.khronos.org/opengl/wiki/Memory_Model#Incoherent_memory_access
- ▶ https://www.khronos.org/opengl/wiki/Synchronization#Implicit_synchronization
- ▶ https://www.khronos.org/opengl/wiki/Buffer_Object_Streaming
- ▶ https://www.khronos.org/opengl/wiki/Buffer_Object
- ▶ <http://webglstats.com/>
- ▶ <https://stackoverflow.com/questions/7954927/passing-a-list-of-values-to-fragment-shader>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>