

Tessellation Shaders

Review

- ▶ We saw vertex and fragment shaders in ETGG 2801/2802
 - ▶ Vertex shader: Takes one vertex as input, outputs one vertex
 - ▶ Fragment shader: Determines color of a pixel

Review

- ▶ We also saw geometry shaders
 - ▶ Take 1-6 vertices as input, output arbitrary number of vertices
 - ▶ Limited in how much *vertex amplification* is feasible

Idea

- ▶ Modern GPU bottleneck: Memory access
 - ▶ Often faster to do computations for one frame, use results, throw them away, recompute on next frame
- ▶ Tessellation shaders are helpful here

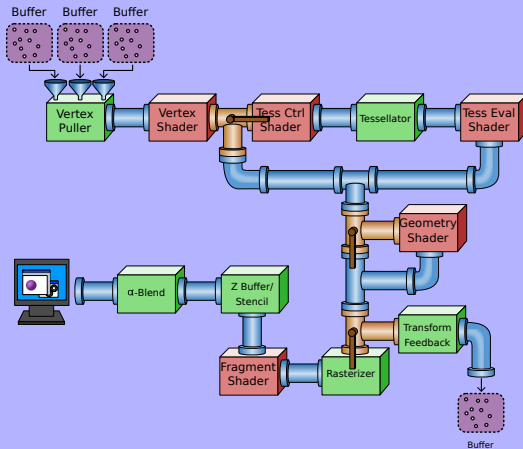
Pipeline

- ▶ Recall the GPU pipeline we saw in 2801/2802:

```
p = []  
for each triangle T:  
    for each vertex v of T:  
        run vertex shader on v to get 2D position p[i]  
for each triple of 2D points in p:  
    compute which pixels are covered by this triangle  
    for each pixel  
        run fragment shader to compute color  
        set pixel to that color
```

Pipeline

- The actual GPU pipeline:



Tessellation Shader

- ▶ Conventional rendering: Input = triangle mesh
 - ▶ Vertices (positions, texture coordinates, normals)
 - ▶ Indices specify which vertices make up triangles
- ▶ Ordinarily, vertices are neither created nor destroyed when drawing meshes
 - ▶ Ignoring the geometry shader...
- ▶ Things are different when using tessellation shader

Input

- ▶ Tessellation shader takes *patch* as input
 - ▶ Patch = Some (constant) number of vertices
 - ▶ Anywhere from 1...32 vertices
- ▶ Example:
`glPatchParameteri(GL_PATCH_VERTICES, numVerts);`
`glDrawElements(GL_PATCHES, count, GL_UNSIGNED_INT, 0);`

Tessellation

- ▶ GPU doesn't attach any geometric meaning to a patch
 - ▶ Just a blob of vertices
- ▶ GPU doesn't consider them “connected” or part of a mesh or anything else
- ▶ Shaders must tell where things end up

VS

- ▶ When using tessellation, VS usually has little to do
 - ▶ Just takes inputs (from buffers) and sends down the pipeline

TCS

- ▶ Tessellation control shader: Runs anywhere from 1...32 times per patch
- ▶ Can see all vertices of the patch at once
- ▶ Produces one output per invocation

Tessellator

- ▶ Tessellator is fixed-function
- ▶ Produces a bunch of vertices
- ▶ Connects together in predetermined pattern to form faces
- ▶ Feeds these to next stage (tessellation evaluation shader)

TES

- ▶ Tessellation evaluation shader: Runs once for each vertex produced by the tessellator
- ▶ Can read all the outputs from the TCS invocations
- ▶ Writes value to `gl_Position`
 - ▶ If no geometry shader, this will be screen space location
 - ▶ If there is GS, we can decide what meaning `gl_Position` will have

GS

- ▶ Recall: GS is run once for each primitive from earlier in pipeline (typically triangles)
- ▶ Can output zero or more triangle strips

FS

- ▶ Recall: Fragment shader: Runs once per pixel
- ▶ Determines what color pixel should have

Example

- ▶ Let's look at a motivating example: Fur/hair generation
- ▶ Idea: We have a mesh; we want to create a shaft of hair at each vertex
 - ▶ Each shaft made up of a line segment

C++ Code

- Our draw code is pretty simple:

```
void draw(){  
    cam->setUniforms();  
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT );  
    Program::setUniform("worldMatrix", mat4::identity() );  
    meshprog->use();  
    M->draw();  
    tessprog->use();  
    glPatchParameteri(GL_PATCH_VERTICES,1);  
    M->draw(GL_PATCHES);  
    fpsText->draw();  
}
```

VS

- ▶ Vertex shader is very simple: Just passes through input data

```
layout(location=0) in vec3 position;  
layout(location=2) in vec3 normal;
```

```
out vec3 v_position;  
out vec3 v_normal;
```

```
void main(){  
    v_position = position;  
    v_normal = normal;  
}
```

TCS

► New stage: Tessellation control shader

```
layout(vertices=1) out;
```

```
in vec3 v_position[];
```

```
in vec3 v_normal[];
```

```
out vec3 tcs_position[];
```

```
out vec3 tcs_normal[];
```

```
void main(){
```

```
    gl_TessLevelOuter[0] = 1;
```

```
    gl_TessLevelOuter[1] = 1;
```

```
    tcs_position[gl_InvocationID] = v_position[gl_InvocationID];
```

```
    tcs_normal[gl_InvocationID] = v_normal[gl_InvocationID];
```

```
}
```

Explanation

- ▶ layout line: Says “For each patch, run TCS this many times”
- ▶ Inputs ('in') are from VS output
 - ▶ Array because we can see all vertices of the patch
- ▶ Outputs will be forwarded to next shader stage
 - ▶ It's an array, but this is a bit of a misnomer
 - ▶ TCS must only write to `tcs_position[gl_InvocationID]` and `tcs_normal[gl_InvocationID]`
 - ▶ Writing to any other slot of an output is an error
- ▶ `gl_TessLevelOuter`: Governs amount of tessellation
 - ▶ `[0]` says how many lines will be generated by tessellator
 - ▶ `[1]` says how many segments there are per line

TES

- ▶ Next we have tessellation evaluation shader:

```
layout(isolines) in;
in vec3 tcs_position[];
in vec3 tcs_normal[];
out float tes_shaftPct;

void main(){
    tes_shaftPct = gl_TessCoord[0];
    vec3 p1 = tcs_position[0] + tes_shaftPct * tcs_normal[0];
    vec4 p = vec4(p1,1.0);
    p = p * worldMatrix;
    p = p * viewMatrix;
    p = p * projMatrix;
    gl_Position = p;
}
```

Explanation

- ▶ layout(isolines) in
 - ▶ Tells tessellator we want it to generate line segments
 - ▶ It's a bit confusing: The TCS and TES work in tandem to direct the tessellator
 - ▶ TCS tells *how much* tessellation to do
 - ▶ TES tells *what format* to use (lines, triangles, quads)
- ▶ Inputs
 - ▶ Arrays because TES can see entire patch's TCS outputs
- ▶ Outputs go to next shader stage (FS in this case)
- ▶ `gl_TessCoord[0]` = How far we are along this line (0...1)
- ▶ `gl_TessCoord[1]` = Which shaft we're doing (0...1)
 - ▶ Since we only create one shaft, this would always be 0
 - ▶ Yes, this seems reversed compared to the tessellation levels' order...

FS

- ▶ FS doesn't show any new concepts:


```
in float tes_shaftPct;
```

```
out vec4 color;
```

```
void main(){  
    color.a = 1.0;  
    color.rgb = vec3(0) + tes_shaftPct * vec3(1);  
}
```

Results

- ▶ What we get:



Results

- ▶ It looks kind of...artificial
- ▶ We'd like the shafts to bend downward like real fur
- ▶ Idea: We do something of a particle simulation
- ▶ Physics:
 - ▶ $s = \int v \, dt$ (where s =distance, v =velocity, t =time)
 - ▶ $v = \int a \, dt$ (where a =acceleration)
- ▶ Since acceleration (due to gravity) is a constant, $v = at$ and $s = \frac{1}{2}at^2$ (ignoring constants)
- ▶ We want to use a quadratic curve

TCS

- ▶ The new TCS: Here, we do constant amount of subdivision:

```
layout(vertices=1) out;
in vec3 v_position[];
in vec3 v_normal[];
out vec3 tcs_position[];
out vec3 tcs_normal[];
void main(){
    gl_TessLevelOuter[0] = 1;
    gl_TessLevelOuter[1] = 8; //<----- new
    tcs_position[gl_InvocationID] = v_position[gl_InvocationID];
    tcs_normal[gl_InvocationID] = v_normal[gl_InvocationID];
}
```



TES

- ▶ We also alter the TES: The constants are chosen from experimentation to give desired appearance

```
layout(isolines) in;  
in vec3 tcs_position[];  
in vec3 tcs_normal[];  
out float tes_shaftPct;  
void main(){  
    tes_shaftPct = gl_TessCoord[0];  
    float t = (tes_shaftPct * 20);  
    vec3 p1 = tcs_position[0] +  
              tes_shaftPct * tcs_normal[0] +  
              vec3(0,-0.005,0)*t*t;  
    vec4 p = vec4(p1,1.0);  
    p = p * worldMatrix; p = p * viewMatrix; p = p * projMatrix;  
    gl_Position = p;  
}
```


Results

- ▶ Gives better appearance than before...But closeups don't look so good



Results

- ▶ If we do more subdivision, hair shafts are smoother (this uses 16 instead of 8)



Idea

- ▶ In practice, we'd measure distance from viewer and subdivide more if object is closer to camera
- ▶ But: Our hair shafts are still pretty far apart
- ▶ Problem: We are requiring each shaft to emanate from a vertex
- ▶ We could get more shafts by adding more vertices
 - ▶ But that's a bad idea (slows rendering)

Procedure

- ▶ First, we'll use entire triangle instead of just one point as input to TCS
- ▶ In draw():
`glPatchParameteri(GL_PATCH_VERTICES,3);`
- ▶ Now, TCS will run once per triangle
- ▶ This would make results sparser

TCS


- Suppose we change the TCS like so:

```
layout(vertices=1) out;
in vec3 v_position[];
in vec3 v_normal[];
out vec3 tcs_position[];
out vec3 tcs_normal[];
void main(){
    gl_TessLevelOuter[0] = 1;
    gl_TessLevelOuter[1] = 8;
    tcs_position[gl_InvocationID] = (v_position[0]+v_position[1]+
        v_position[2])/3.0;
    tcs_normal[gl_InvocationID] = (v_normal[0]+v_normal[1]+
        v_normal[2])/3.0;
}
```

- Now, each triangle gets one shaft at its centroid

Results

- ▶ Screenshot



Multiple Hairs

- ▶ How to create more hairs?
- ▶ Want number of hairs to be proportional to triangle area
 - ▶ Can get this via cross product
- ▶ Once we have number of hairs, how can we distribute randomly?

Barycentric

- ▶ Recall: Barycentric coordinates: For any u, v, w The point $v[0]*u + v[1]*v + v[2]*w$ will be within the triangle if:
 - ▶ All three u, v, w are between 0 and 1
 - ▶ $u+v+w = 1$
- ▶ u, v, w are barycentric coordinates of triangle

Randomness

- ▶ Next issue: How can we obtain random numbers on GPU?
- ▶ On CPU, we use `rand()`
- ▶ How does `rand()` work?

LCG

- ▶ Most systems use *linear congruential generator* for rand()
 - ▶ Input: A seed
 - ▶ Compute: $\text{seed} = (a * \text{seed}) \bmod b$
 - ▶ Return seed as random number
- ▶ Notice: Generator continually updates seed

Problem

- ▶ What do we choose for a and b ?
 - ▶ Lots of mathematical theory
 - ▶ Recommended values: $a=0x7ffffff$, $b=16807$
- ▶ Problem: Computing $a * \text{seed}$ may overflow
 - ▶ We can address this by some careful coding
 - ▶ But: Can we avoid the problem entirely?

Xorshift

- ▶ Another RNG that does not have this problem is Xorshift
- ▶ Basic structure: We have some global variable named “state”
 - ▶ `state = state ^ (state << 13);`
 - ▶ `state = state ^ (state >> 17);`
 - ▶ `state = state ^ (state << 15);`
- ▶ Often we want a value between 0...1, so we convert this uint to a float:
`randomValue = float(state & 0x7fffffff) / float(0x7fffffff);`

Code

- ▶ We can package this into a function:

```
uint state;  
float xorshift(){  
    state = state ^ (state << 13);  
    state = state ^ (state >> 17);  
    state = state ^ (state << 15);  
    return float(state & 0x7fffffff) / float(0x7fffffff);  
}
```

- ▶ Our only problem...How to initialize state?

Initialize

- ▶ Shader invocations don't share global variables
- ▶ So we need to initialize state to different value in each shader
- ▶ We have a primitive number (patch ID)
- ▶ And we know our shaft number:

```
uint shaftNumber = uint( gl_TessCoord[1] * gl_TessLevelOuter[0] );
```

 - ▶ `gl_TessCoord[1]` is value in range `[0,1)`: 0 for first shaft
 - ▶ `gl_TessLevelOuter[0]` tells how many shafts we have total
 - ▶ So we multiply to get integer value 0,1,2,...

Initialization

- ▶ Now we can write an initialization function: Stir primitive ID and shaft number together

```
void xorshiftInit(uint v){  
    state = gl_PrimitiveID ^ (v<<16);  
}
```

- ▶ Call it as:
`xorshiftInit(shaftNumber);`

Putting it Together

- ▶ C code doesn't change:

```
cam->setUniforms();  
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT );  
Program::setUniform("worldMatrix", mat4::identity() );  
meshprog->use();  
M->draw();  
tessprog->use();  
glPatchParameteri(GL_PATCH_VERTICES,3);  
M->draw(GL_PATCHES);  
fpsText->draw();
```

Putting it Together

► TCS:

```
layout(vertices=3) out;  
in vec3 v_position[];  
in vec3 v_normal[];  
out vec3 tcs_position[];  
out vec3 tcs_normal[];  
void main(){  
    gl_TessLevelOuter[0] = 16; //16 shafts per triangle  
    gl_TessLevelOuter[1] = 8;  //8 segments per shaft  
    tcs_position[gl_InvocationID] = v_position[gl_InvocationID];  
    tcs_normal[gl_InvocationID] = v_normal[gl_InvocationID];  
}
```

Note

- ▶ We should scale number of shafts per triangle based on triangle area...
 - ▶ We'll overlook that for now
- ▶ Notice TCS runs 3 times per patch
- ▶ We've set patch size to 3 (C code)

Putting it Together

► TES:

```
layout(isolines) in;  
in vec3 tcs_position[];  
in vec3 tcs_normal[];  
out float tes_shaftPct;
```

```
void main(){  
    ...  
}
```

TES

- ▶ What does TES need to do?
- ▶ First: Initialize RNG

```
tes_shaftPct = gl_TessCoord[0];  
uint shaftNumber = uint(  
    gl_TessCoord[1] * gl_TessLevelOuter[0] );  
xorshiftInit(shaftNumber);
```

TES

- ▶ Next, generate random barycentric coordinates

```
float u = xorshift();  
float v = xorshift();  
if( u+v > 1.0 ){  
    u = 1-u;  
    v = 1-v;  
}  
float w = 1.0-u-v;
```

- ▶ Now, $u+v+w=1$ and all of u,v,w are between 0 and 1

TES

- ▶ We can write a small barycentric interpolation routine:

```
vec3 interpolate(float u, float v, float w,  
                vec3 a, vec3 b, vec3 c ){  
    vec3 p0 = u * a;  
    vec3 p1 = v * b;  
    vec3 p2 = w * c;  
    return (p0+p1+p2);  
}
```

- ▶ If this is given triangle vertices or normals, it does the required interpolation

TES

- Compute position and normal where the hair shaft is:

```
vec3 basePos = interpolate(  
    u,v,w,  
    tcs_position[0],tcs_position[1],tcs_position[2]);
```

```
vec3 N = interpolate(  
    u,v,w,  
    tcs_normal[0],tcs_normal[1],tcs_normal[2]);
```

TES


- ▶ Last thing: Compute output for rasterizer:

```
float t = (tes_shaftPct * 20);  
vec3 p1 = basePos + tes_shaftPct * 0.35 * N + vec3(0,-0.0005,0)*t  
    *t;  
vec4 p = vec4(p1,1.0);  
p = p * worldMatrix;  
p = p * viewMatrix;  
p = p * projMatrix;  
gl_Position = p;
```

- ▶ The 0.35 governs how long the shaft is
- ▶ The -0.0005 governs the amount of vertical bend to the shaft

Result

- ▶ Our result:



- ▶ Hmmmm... It looks kind of...patchy

Problem

- ▶ Random number generators (LCG's or Xorshift) are good for sequences of numbers
- ▶ But if the seeds are not random, they tend to produce correlated sequences
- ▶ One solution: We can hash the input seed for the generator

Hash

- ▶ One good hash: A Wang hash (developed by Thomas Wang)

```
uint hash(uint v){  
    v = 9*( (v >> 16 ) ^ (v ^ 61));  
    v ^= (v>>4);  
    v = v * 668265261;  
    return v ^ (v>>15);  
}
```

- ▶ Use it:

```
void xorshiftInit(uint v){  
    state = gl_PrimitiveID ^ (v<<16);  
    state = hash(state)  
}
```


Results

- ▶ Much better




Note

- ▶ If we are simulating scene where we mostly look at shafts edge-on (ex: grass), we can use transparency to help appearance
 - ▶ When drawing: `glDepthMask(0)` before rendering hair; `glDepthMask(1)` afterwards
 - ▶ In FS: `out.a = 1.0 - tes_shaftPos`



Final Render

- ▶ With lighting



Sources

- ▶ <https://www.khronos.org/opengl/wiki/Tessellation>
- ▶ https://www.khronos.org/opengl/wiki/Tessellation_Control_Shader
- ▶ https://www.khronos.org/opengl/wiki/Tessellation#Tessellation_primitive_generation
- ▶ Kostas Anagnostou. Rendering Fur Using Tessellation.
<https://interplayoflight.wordpress.com/2014/12/31/rendering-fur-using-tessellation/>
- ▶ Nathan Reed. Quick and Easy GPU Random Numbers in D3D11.
<http://www.reedbeta.com/blog/quick-and-easy-gpu-random-numbers-in-d3d11/>
- ▶ Integer Hashing. <http://www.burtleburtle.net/bob/hash/integer.html>
- ▶ Adam Swaab. Random Point in a Triangle - Barycentric Coordinates.
<https://adamswaab.wordpress.com/2009/12/11/random-point-in-a-triangle-barycentric-coordinates/>
- ▶ Random Numbers. http://www.eternallyconfuzzled.com/tuts/algorithms/js/tut_rand.aspx
- ▶ George Marsaglia. Xorshift RNGs. <http://www.jstatsoft.org/v08/i14/paper>
- ▶ Humus. Phone-Wire AA. <http://www.humus.name/index.php?page=3D&ID=89>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>