

SSE

Motivation

- ▶ Often, we have repetitive operations to be executed on several data items
- ▶ We can speed things up by doing them in parallel
- ▶ Concept: SIMD: Single Instruction Multiple Data
- ▶ To take advantage of this, we might need to reorganize our code/data structures

Background

- ▶ We'll only consider x86 architecture here
 - ▶ ARM has similar instructions (Neon)
 - ▶ Other platforms have SIMD instructions too
- ▶ Basic idea:
 - ▶ Load several data items into a single register
 - ▶ Do one operation on all the data items in parallel
 - ▶ (Usually) store the data items back to RAM

Review

- ▶ Recall x86 CPU architecture...
 - ▶ Registers: Temporary storage locations
 - ▶ Used by CPU to hold data to be processed
 - ▶ 32 bit CPU: Seven GPR's: eax, ebx, ecx, edx, esi, edi, ebp
 - ▶ 64 bit CPU: Sixteen GPR's: rax, rbx, rcx, rdx, rsi, rdi, rbp, r8, r9, r10, r11, r12, r13, r14, r15

FP

- ▶ Floating point values are stored differently from integers
 - ▶ Like scientific notation
 - ▶ Ex: $3014.159 = 3.014159 \times 10^3$
- ▶ Since computers are binary, we use powers of 2, not powers of 10
 - ▶ Ex: $1.011011_2 \times 2^4$
- ▶ Float: 32 bits
 - ▶ 23 bits for mantissa, 8 bits for exponent, 1 bit for sign
- ▶ Double: 64 bits
 - ▶ 52 bits for mantissa, 11 bits for exponent, 1 bit for sign
- ▶ Extended: 80 bits
 - ▶ 62 bits for mantissa, 15 bits for exponent, 1 bit for sign, 1 bit for integer flag

FP

- ▶ x86 has special unit (FPU) to handle floating point operations
- ▶ Since FP numbers won't fit in 32 bit registers, Intel added 8 80-bit registers for FP values (way back in early 80's)
- ▶ Plus separate set of instructions to process them

Prehistory

- ▶ Early SIMD instruction set (1996): MMX
 - ▶ Officially, it doesn't stand for anything
 - ▶ You can't trademark an acronym!
 - ▶ But it could stand for "MultiMedia eXtensions"
- ▶ Used part of FP registers to operate on 64 bits at a time
 - ▶ CPU was only 32 bits, so couldn't use GPR's
- ▶ Limited, clunky, slow
- ▶ We ignore this

Stone Age

- ▶ Intel's next try: SSE (Pentium III, 1999)
- ▶ SSE stands for “Streaming SIMD”
 - ▶ Apparently, Intel gave up on the whole trademark thing
- ▶ A set of new instructions + 8 new registers
 - ▶ 128 bits wide
 - ▶ xmm0 - xmm7

Bronze Age

- ▶ SSE2: Pentium 4, 2001
- ▶ Additional SSE instructions
- ▶ Still uses xmm registers

Iron Age

- ▶ SSE3: Pentium 4, 2004
- ▶ More SSE instructions
- ▶ Same xmm registers

Medieval Period

- ▶ SSSE3: Core 2
- ▶ Additional SSE instructions
 - ▶ Super Streaming SIMD Extensions?
- ▶ Same xmm registers

Victorian Era

- ▶ SSE4: Intel Core Penryn/Nehalem, 2007
- ▶ Yet more SSE instructions
- ▶ Same xmm registers, but 64 bit code has 8 more of them (xmm8-xmm15)

Modern Age

- ▶ AVX: Intel Core, 2011
- ▶ New 256 bit registers (YMM0-YMM15)
 - ▶ Note: XMM## = Lower half of YMM##
- ▶ Plus some instructions for these registers
- ▶ And some three-operand instructions (finally!)

The Future

- ▶ AVX-512: Mid-2017
- ▶ Adds 512 bit ZMM registers
- ▶ Not widely available
 - ▶ I.e., I don't have a machine that supports these
 - ▶ So we ignore them

Code

- ▶ We'll focus on SSE/AVX/AVX2 since support is fairly widespread
- ▶ And they are reasonably easy to work with
 - ▶ Especially compared to the floating point registers
 - ▶ MMX has some snags that make it more difficult to mix with other code

Data Types

- ▶ XMM registers are 128 bits = 16 bytes
- ▶ They can be treated as:
 - ▶ 16 single-byte quantities
 - ▶ 8 shorts
 - ▶ 4 ints
 - ▶ 4 floats
 - ▶ 2 longs
 - ▶ 2 doubles

Data Types

- ▶ YMM registers are 256 bits = 32 bytes
- ▶ They can be treated as:
 - ▶ 32 single-bytes
 - ▶ 16 shorts
 - ▶ 8 ints
 - ▶ 8 floats
 - ▶ 4 longs
 - ▶ 4 doubles

Intrinsics

- ▶ Too difficult to use assembly language directly
- ▶ Compilers provide alternatives: *intrinsics*
- ▶ They look like ordinary functions, but they translate directly into assembly instructions

Using

- ▶ GCC/Intel compilers:
 - ▶ `#include <mmintrin.h> //mmx`
 - ▶ `#include <xmmintrin.h> //sse`
 - ▶ `#include <emmintrin.h> //sse2`
 - ▶ `#include <pmmmintrin.h> //sse3`
 - ▶ `#include <tmmintrin.h> //ssse3`
 - ▶ `#include <smmmintrin.h> //sse4.1`
 - ▶ `#include <nmmmintrin.h> //sse4.2`
 - ▶ `#include <immintrin.h> //avx, avx2`
- ▶ In Visual Studio, most are in `intrin.h`, but a few are in `immintrin.h`

Compile

- ▶ gcc
 - ▶ Compile with `-mfpmath=sse -mmmx -msse -msse2 -msse3 -msse4.1 -msse4.2 -msse4 -mavx -mavx2`
- ▶ Visual Studio
 - ▶ I suppose there's some options somewhere for this...

Note

- ▶ See <https://software.intel.com/sites/landingpage/IntrinsicsGuide/#> for the full reference
- ▶ Bookmark this and put a copy under your pillow
- ▶ You will *need* this site!

Data Types

- ▶ 128-bit data types: `__m128`, `__m128i`, `__m128d`
 - ▶ `m128i` = Integers
 - ▶ `m128d` = floating point
 - ▶ They're the same bit-wise, but system can optimize if it knows what type's expected to go in the register
- ▶ 256-bit data types: `__m256`, `__m256i`, `__m256d`

Functions

- ▶ Naming pattern: **_mm_operation_type**
 - ▶ operation is the operation being performed (add, subtract, etc.)
 - ▶ type tells what kind of data
 - ▶ ps = packed floats (4 at a time)
 - ▶ pd = packed doubles (2 at a time)
 - ▶ si128 = integer values (byte, short, int)

Examples

- ▶ We'll start off with the float variants
- ▶ Four floats per xmm register

load_ps

- ▶ Load parallel (packed?) single precision
 - ▶ `float* p = ...`
 - ▶ `__m128 v = _mm_load_ps(p)`
 - ▶ `v[0] = *p`
 - ▶ `v[1] = *(p+1)`
 - ▶ `v[2] = *(p+2)`
 - ▶ `v[3] = *(p+3)`
- ▶ Important: Address pointed to by `p` must be on 16 byte boundary

Alignment

- ▶ Note: Some load instructions require data to be 16-byte aligned
- ▶ If it's not: Processor will issue a fault
- ▶ Ex: This may or may not work:
`float F[4] = { ... };`
`__m128 v = _mm_load_ps(F);`
- ▶ C++ 11 provides an `alignas` operator we can use:
`alignas(16) float F[4] = {...};`
`__m128 v = _mm_load_ps(F);`

Problem

- ▶ Dynamically allocated memory (malloc, new) is not guaranteed to be aligned
- ▶ Likewise for memory in vectors
- ▶ So this has a chance of crashing:
float* F = new float[n];
...fill it up...
__m128 v = _mm_load_ps(F);
- ▶ Likewise:
vector<float> F(n);
...fill it up...
__m128 v = _mm_load_ps(F.data());

Solution

- ▶ We can define our own type:
`struct alignas(16) XMMFloat{
 float v[4];
};`
- ▶ Then we can use new and vectors properly:
`XMMFloat* F = new XMMFloat[n];
...fill it up...
__m128 v = _mm_load_ps((float*)F);`
- ▶ Likewise:
`vector<XMMFloat> F(n);
...fill it up...
__m128 v = _mm_load_ps((float*)F.data());`

load_ps1

- ▶ Load and copy one single precision item
 - ▶ `float* p = ...`
 - ▶ `__m128 v = _mm_load1_ps(p)`
 - ▶ `v[0] = *p`
 - ▶ `v[1] = *p`
 - ▶ `v[2] = *p`
 - ▶ `v[3] = *p`
- ▶ Notice: It replicates the same thing to four places
- ▶ The documentation isn't clear if `p` must be aligned

load_ss

- ▶ Load scalar single precision: Loads one item and zeros the rest
 - ▶ `float* p = ...`
 - ▶ `__m128 v = _mm_load_ss(p)`
 - ▶ `v[0] = *p`
 - ▶ `v[1] = 0`
 - ▶ `v[2] = 0`
 - ▶ `v[3] = 0`
- ▶ No required alignment

loadr_ps

- ▶ Load reversed parallel (packed) single precision
- ▶ `__m128 v = _mm_loadr_ps(p)`
 - ▶ `v[3] = *p`
 - ▶ `v[2] = *(p+1)`
 - ▶ `v[1] = *(p+2)`
 - ▶ `v[0] = *(p+3)`
- ▶ Must be aligned

loadu_ps

- ▶ Load unaligned packed single precision
- ▶ `__m128 v = _mm_loadu_ps(p)`
 - ▶ `v[0] = *p`
 - ▶ `v[1] = *(p+1)`
 - ▶ `v[2] = *(p+2)`
 - ▶ `v[3] = *(p+3)`
- ▶ Same as `load_ps`, but `p` need not be on a 16-byte aligned address
- ▶ Older CPU's had speed penalty for unaligned; modern ones don't (for *most* accesses)

Store

- ▶ `v = __m128` value; `p=`float pointer
- ▶ `_mm_store_ps(p, v)`
 - ▶ Stores contents of `v` to `p[0]...p[3]`
- ▶ `_mm_store_ps1(p, v)`
 - ▶ Replicates `v[0]` four times
- ▶ `_mm_storer_ps(p, v)`
 - ▶ Stores in reverse order
- ▶ `_mm_storeu_ps(p, v)`
 - ▶ Unaligned store (maybe slower)

Example

- ▶ Example code: [ex0.cpp](#)

- ▶ Output:

3 1 4 5

0 0 0 0

Computation

- ▶ OK, so now we know how to load and store data.
- ▶ What can we do with it?

Arithmetic

- ▶ Add 4 floats in parallel: `__m128 v3 = _mm_add_ps(v1,v2)`
- ▶ Subtract: `v3 = _mm_sub_ps(v1,v2)`
- ▶ Multiply: `v3 = _mm_mul_ps(v1,v2)`
- ▶ Divide: `v3 = _mm_div_ps(v1,v2)`
- ▶ GCC provides overloads for `+`, `-`, `*`, and `/` as well as unary negation
 - ▶ VS might do the same thing...

Example

- ▶ Example: [ex1.cpp](#)
- ▶ Output:
12 3 10 10

Arithmetic

- ▶ Reciprocal: `__m128 v2 = _mm_rcp_ps(v1)`
- ▶ Reciprocal square root: `v2 = _mm_rsqrt_ps(v1)`
- ▶ Square root: `v2 = _mm_sqrt_ps(v1)`
- ▶ Zero: `v = _mm_setzero_ps()`
 - ▶ Zero out the whole register
- ▶ Minimum: `v3 = _mm_min_ps(v1,v2)`
- ▶ Maximum: `v3 = _mm_max_ps(v1,v2)`
- ▶ Ceiling: `v2 = _mm_ceil_ps(v1)`
- ▶ Floor: `v2 = _mm_floor_ps(v1)`
- ▶ Round: `v2 = _mm_round_ps(v1, mode)`
 - ▶ Mode = `_MM_FROUND_TO_NEAREST_INT` or `_MM_FROUND_TO_ZERO`

Example

- ▶ Let's see a real-life example: Scaling loudness in a WAV file
- ▶ Explain: WAV files: Header followed by data
- ▶ Data format: Several possible...

PCM

- ▶ PCM: Easiest to understand
- ▶ Format code = 1 (integer) or 3 (float)
- ▶ Integer format: 8, 16, 24, or 32 bits per sample
 - ▶ 8 bit = unsigned, 16/24/32 bit = signed
- ▶ Float format: 32 bits per sample, -1...1 range
- ▶ Directly measures amplitude of signal
- ▶ Usually one or two channels

ADPCM

- ▶ ADPCM: Provides compression
- ▶ Input: 16 bit PCM value
- ▶ Output: 4 bit quantized value: The delta between previous sample and this one
- ▶ As long as signal doesn't change too rapidly, ADPCM gives good reproduction at 25% of the bitrate

μ -law

- ▶ Quantizes 14 bit samples to 8 bit
 - ▶ Gives increased resolution in critical areas of waveform without requiring larger sample size
- ▶ For range from -8192...8191:
 - ▶ Break range into 18 intervals of different sizes
 - ▶ Closer interval is to zero, smaller it is
 - ▶ Further interval is from zero, larger it is
 - ▶ Divide each interval into 16 pieces
- ▶ For each 14 bit input sample:
 - ▶ Determine which interval/piece it falls in
 - ▶ Output 8 bit value: 4 bits for interval identifier, 4 for piece identifier
- ▶ Result: Input parts near 0 get encoded with more accuracy; further away get less accuracy
- ▶ Works well for speech; used in analog telephones

Code

- ▶ We'll assume PCM/float format
- ▶ Header to load wave files: [Wave.h](#)
- ▶ Stopwatch for timing: [Stopwatch.h](#)
- ▶ Non-SSE program to adjust loudness: [loudness.cpp](#)
- ▶ SSE program to adjust: [loudnessxmm.cpp](#)

Results

- ▶ Input: WAV file, about 1 minute in length
- ▶ Machine: Core i7
- ▶ Average of five executions (microseconds):

Non-SIMD	SIMD
4439	2529

Assignment

- ▶ Write a program which takes three command line parameters.
 - ▶ The first is the name of a WAV file encoded using the floating point format.
 - ▶ The second is the echo decay (a number from 0 to 1).
 - ▶ The third is the echo delay.
 - ▶ Using these parameters, add echo using SIMD intrinsics.
 - ▶ Write the output to the file "out.wav".
 - ▶ Report the total time required for the operation.
- ▶ There are example files on the class website.

Sources

- ▶ <https://software.intel.com/en-us/articles/intel-software-development-emulator#faq>
- ▶ <http://www.tomshardware.com/reviews/intel-drops-pentium-brand,1832-2.html>
- ▶ <https://www.mathworks.com/matlabcentral/answers/93455-what-is-the-sse2-instruction-set-how-can-i-check-to-see-if-my-processor-supports-it?requestedDomain=www.mathworks.com>
- ▶ <http://softpixel.com/~cwright/programming/simd/ssse3.php>
- ▶ <https://software.intel.com/sites/default/files/m/8/b/8/D9156103.pdf>
- ▶ [https://msdn.microsoft.com/en-us/library/y08s279d\(v=vs.100\).aspx](https://msdn.microsoft.com/en-us/library/y08s279d(v=vs.100).aspx)
- ▶ <http://www.linuxjournal.com/content/introduction-gcc-compiler-intrinsics-vector-processing?page=0,2>
- ▶ <http://en.cppreference.com/w/cpp/language/alignas>
- ▶ <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>
- ▶ MSDN SSE reference
- ▶ <http://ok-cleek.com/blogs/?p=20540>
- ▶ <http://www.liranuna.com/sse-intrinsics-optimizations-in-popular-compilers/>
- ▶ <https://www.cilkplus.org/tutorial-pragma-simd>
- ▶ <https://stackoverflow.com/questions/2804902/whats-the-difference-between-logical-sse-intrinsics>
- ▶ <https://en.cppreference.com/w/cpp/language/alignas>
- ▶ <http://www.ent.mrt.ac.lk/rds/index.php/curriculum-support/dsd-projects/30-batch-07-dsd-projects/57-wav-to-adpcm-converter>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>