

# Swizzling

# Motivation

- ▶ SSE intrinsics and datatypes
- ▶ We've seen arithmetic operations
- ▶ Now we'll look at *swizzling*

# Swizzle

- ▶ Often, we need to adjust data layout before further processing
- ▶ SSE provides a number of data swizzle instructions
- ▶ But first, we need to examine details of data layout

## Little Endian

- ▶ Intel CPU's are little endian
- ▶ Means LSB stored first in RAM
- ▶ Intel followed this convention for XMM registers too
  - ▶ First float in RAM goes to low slot of XMM register
- ▶ This might be inconvenient in some cases (ex: for matrices)
  - ▶ In that case use `_mm_loadr_ps`

## Illustration

- ▶ Suppose we have this:  
`alignas(16) float F[] = {1,2,3,4};`  
`__m128 v = _mm_load_ps(F);`
- ▶ Question: Does the lowest slot of v get 1 or 4?
  - ▶ Does it matter?

## Illustration

- ▶ We might not care!

- ▶ Consider:

```
__m128 v = _mm_load_ps(F);  
__m128 w = _mm_set1_ps(42.0);  
v = _mm_add_ps(v,w);  
_mm_store_ps(F,v);
```

- ▶ Doesn't matter which end of v gets F[0]
  - ▶ Store will be consistent with load
  - ▶ That's all we need

## Endian

- ▶ Let's take a look at an example instruction where it does make a difference
- ▶ Shuffle packed single precision
- ▶ `__m128 v = _mm_shuffle_ps(__m128 a, __m128 b, uint8_t s )`
  - ▶ `s` is 8 bits. Call the bits DDCCBBAA
  - ▶ `v[0] = a[AA]`
  - ▶ `v[1] = a[BB]`
  - ▶ `v[2] = b[CC]`
  - ▶ `v[3] = b[DD]`

## Example

- ▶ Let's go back to our previous example and add a shuffle:

```
alignas(16) float F[] = {1,2,3,4};  
__m128 v = _mm_load_ps(F);  
//broadcast slot 0 to all outputs  
__m128 x = _mm_shuffle_ps(v,v,0);  
_mm_store_ps(F,x);  
for(int i=0;i<4;++i)  
    cout << F[i] << " ";
```

- ▶ This outputs: 1 1 1 1
- ▶ So the *first* thing in RAM goes to slot 0 of the XMM register



# Shuffle

- ▶ Shuffle for doubles: `v = _mm_shuffle_pd(__m128d a, __m128d b, uint8_t s)`
  - ▶ If low bit of `s` is zero: Copy `a[0]` to `v[0]`, else `a[1]` to `v[0]`
  - ▶ If second bit of `s` is zero: copy `b[0]` to `v[1]`, else `b[1]` to `v[1]`
- ▶ Shuffle for ints: `v = _mm_shuffle_epi32(__m128i a, __m128i b, int s)`
  - ▶ Shuffles 32 bit chunks. Essentially the same as `_mm_shuffle_ps`

# Shuffle


- ▶ Shuffle for bytes is more challenging because 16 bytes fit in one XMM register
- ▶ So an xmm argument is used instead
- ▶ `__m128i v = _mm_shuffle_epi8( __m128i a, __m128i b )`
  - ▶ b is treated as 16 single-byte elements
- ▶ For all 16 elements (i.e., let  $i=0\dots15$ ):
  - ▶ If highest bit of  $b[i]$  is one:  $v[i] = 0$
  - ▶ Else:  $v[i] = a[ b[i] \& 0xf ]$
- ▶ We can use this for shorts as well: Just set values of b to move adjacent elements around correctly

## Note

- ▶ Handy function: Set xmm register to constant:  
`__m128i v = _mm_set_epi8( uint8_t slot15, uint8_t slot14, ... ,  
uint8_t slot1, uint8_t slot0 )`
- ▶ Also `set_epi16`, `set_epi32`, `set_epi64x`
  - ▶ There's a `set_epi64`, but it uses `__m64` type instead of `__int64` type

# Unpack


- ▶ Unpack:  $v = \text{\_mm\_unpackhi\_ps}(a,b)$ 
  - ▶  $v[0] = a[2]$
  - ▶  $v[1] = b[2]$
  - ▶  $v[2] = a[3]$
  - ▶  $v[3] = b[3]$



- ▶ More restricted version of shuffle

# Unpack

- ▶ Unpack:  $v = \text{\_mm\_unpacklo\_ps}(a,b)$ 
  - ▶  $v[0] = a[0]$
  - ▶  $v[1] = b[0]$
  - ▶  $v[2] = a[1]$
  - ▶  $v[3] = b[1]$




# Unpack

- ▶ There's one for bytes: `v = _mm_unpacklo_epi8(a,b)`
  - ▶ `v[0]=a[0]`
  - ▶ `v[1]=b[0]`
  - ▶ `v[2]=a[1]`
  - ▶ `v[3]=b[1]`
  - ▶ `v[4]=a[2]`
  - ▶ ...
  - ▶ `v[15]=b[7]`
- ▶ And `_mm_unpackhi_epi8(a,b)`
  - ▶ Same idea, but uses `a[8...15]` and `b[8...15]`
- ▶ Similar instructions for 16, 32, 64 bit ints


# Replicate

- ▶ Replicate even items of XMM register
- ▶ `__m128 v = _mm_moveldup_ps(a)`
  - ▶  $v[0] = a[0]$
  - ▶  $v[1] = a[0]$
  - ▶  $v[2] = a[2]$
  - ▶  $v[3] = a[2]$



# Replicate

- ▶ For odd items
- ▶ `__m128 v = _mm_movehdup_ps(a)`
  - ▶  $v[0] = a[1]$
  - ▶  $v[1] = a[1]$
  - ▶  $v[2] = a[3]$
  - ▶  $v[3] = a[3]$





## Move

- ▶ Copy top two items of a to top two slots of result; copy top two items of b to bottom two slots of result
- ▶ `__m128 v = _mm_movehl_ps( a, b )`
  - ▶ `v[0] = b[2]`
  - ▶ `v[1] = b[3]`
  - ▶ `v[2] = a[2]`
  - ▶ `v[3] = a[3]`

## Move

- ▶ Copy bottom two items of a to bottom two items of result; copy bottom two items of b to top two slots of result
- ▶ `__m128 v = _mm_movelh_ps( a, b )`
  - ▶ `v[0] = a[0]`
  - ▶ `v[1] = a[1]`
  - ▶ `v[2] = b[0]`
  - ▶ `v[3] = b[1]`

## SSE4

- ▶ `__m128 v = _mm_blend_ps( __m128 a, __m128 b, int s )`
  - ▶ `v[0] = ( s & 1 ) ? b[0] : a[0]`
  - ▶ `v[1] = ( s & 2 ) ? b[1] : a[1]`
  - ▶ `v[2] = ( s & 4 ) ? b[2] : a[2]`
  - ▶ `v[3] = ( s & 8 ) ? b[3] : a[3]`
- ▶ Likewise for `_mm_blend_epi16( __m128i a, __m128i b, int s )`
  - ▶ 16 bit integers; uses low 8 bits of s
- ▶ And `_mm_blend_epi32( __m128i a, __m128i b, int s )`
  - ▶ Uses low four bits of s
- ▶ And `_mm_blend_pd` (for doubles)
  - ▶ Uses two bits from s
- ▶ Similar functions for AVX, but twice as many items in arguments

## Blend

- ▶ For blending byte quantities, function is a bit different
- ▶ `__m128i v = _mm_blendv_epi8(__m128i a, __m128i b, __m128i c)`
  - ▶  $v[i] = (c[i]) ? b[i] : a[i]$ 
    - ▶ For  $i = 0 \dots 15$
- ▶ `__m256i v = _mm256_blendv_epi8( __m256i a, __m256i b, __m256i c)`
  - ▶ Same thing but for 32 slots

## Example

- ▶ Let's see a real-world example of swizzling
- ▶ Suppose we want to do matrix-vector multiplication
- ▶ Suppose we have vector stored as `__m128`
- ▶ Our matrix is stored as four `__m128`'s (rows)
- ▶ Question: How to do multiplication?
  - ▶ (Work out in class)

## Problem

- ▶ Not obvious how to do it
- ▶ The data we need is spread out among four `__m128`'s
- ▶ Postmultiplying vector would be easy:
  - ▶ Dot products of matrix rows with `vec4`.
  - ▶ Done!

## Problem

- ▶ If we stored matrix as *column major* it would be easy to do vector-matrix multiply
- ▶ But: C/C++ defaults to row-major order
  - ▶ Changing to column major might involve lots of data shuffling at load/store time
- ▶ And: We can't do matrix-vector multiply easily if we store data as column major

# Swizzle

- ▶ This is where swizzling can help
- ▶ When we want to multiply:
  - ▶ Compute transpose of matrix
  - ▶ Then do dot product of vector with rows
- ▶ Computing transpose is a frequent operation, so we'd like to be able to do it anyway
- ▶ How to do?
  - ▶ (Work out in class)



# Transpose

- ▶ Suppose we have four matrix rows in r1, r2, r3, r4
- ▶ Declare temporaries w,x,y,z
- ▶ Outputs: tr1,tr2,tr3,tr4: The transposed rows
- ▶ How to do?

r1	1,2,3,4	w		tr1	
r2	5,6,7,8	x		tr2	
r3	9,10,11,12	y		tr3	
r4	13,14,15,16	z		tr4	

# Transpose

► `w = _mm_unpackhi_ps( r1, r2 );`

r1	1,2,3,4	w	3,7,4,8	tr1	
r2	5,6,7,8	x		tr2	
r3	9,10,11,12	y		tr3	
r4	13,14,15,16	z		tr4	

## Test

► `x = _mm_unpackhi_ps( r3, r4 );`

r1	1,2,3,4	w	3,7,4,8	tr1	
r2	5,6,7,8	x	11,15,12,16	tr2	
r3	9,10,11,12	y		tr3	
r4	13,14,15,16	z		tr4	

# Transpose

► `y = _mm_unpacklo_ps( r1, r2 );`

r1	1,2,3,4	w	3,7,4,8	tr1	
r2	5,6,7,8	x	11,15,12,16	tr2	
r3	9,10,11,12	y	1,5,2,6	tr3	
r4	13,14,15,16	z		tr4	

# Transpose

► `z = _mm_unpacklo_ps( r3, r4 );`

r1	1,2,3,4	w	3,7,4,8	tr1	
r2	5,6,7,8	x	11,15,12,16	tr2	
r3	9,10,11,12	y	1,5,2,6	tr3	
r4	13,14,15,16	z	9,13,10,14	tr4	

# Transpose

► `tr1 = _mm_movelh_ps( y , z );`

r1	1,2,3,4	w	3,7,4,8	tr1	1,5,9,13
r2	5,6,7,8	x	11,15,12,16	tr2	
r3	9,10,11,12	y	1,5,2,6	tr3	
r4	13,14,15,16	z	9,13,10,14	tr4	

# Transpose

► `tr2 = _mm_movehl_ps( z , y );`

r1	1,2,3,4	w	3,7,4,8	tr1	1,5,9,13
r2	5,6,7,8	x	11,15,12,16	tr2	2,6,10,14
r3	9,10,11,12	y	1,5,2,6	tr3	
r4	13,14,15,16	z	9,13,10,14	tr4	

# Transpose

► `tr3 = _mm_movelh_ps( w , x );`

r1	1,2,3,4	w	3,7,4,8	tr1	1,5,9,13
r2	5,6,7,8	x	11,15,12,16	tr2	2,6,10,14
r3	9,10,11,12	y	1,5,2,6	tr3	3,7,11,15
r4	13,14,15,16	z	9,13,10,14	tr4	



# Transpose

► `trow4 = _mm_movehl_ps( x , w );`

r1	1,2,3,4	w	3,7,4,8	tr1	1,5,9,13
r2	5,6,7,8	x	11,15,12,16	tr2	2,6,10,14
r3	9,10,11,12	y	1,5,2,6	tr3	3,7,11,15
r4	13,14,15,16	z	9,13,10,14	tr4	4,8,12,16

Done!

- ▶ Now, tr1,tr2,tr3,tr4 are the transposed matrix's rows
- ▶ We can now use the dot product intrinsic
- ▶ Recall: `_mm_dp_ps(a,b,f)`
  - ▶ `f` = scalar: Flags
  - ▶ For vec4/vec4 dot product replicated to all four slots of output: `f=0xff`

## Compute

- ▶ `dpx = _mm_dp_ps( vec, tr1, 0xf1 );`
- ▶ `dpy = _mm_dp_ps( vec, tr2, 0xf2 );`
- ▶ `dpz = _mm_dp_ps( vec, tr3, 0xf4 );`
- ▶ `dpw = _mm_dp_ps( vec, tr4, 0xf8 );`
- ▶ These are x,y,z,w of result

## Store

- ▶ We need to combine these four items into one xmm register
- ▶ We can do this with bitwise-or
  - ▶ `dpx = _mm_or_ps( dpx,dpy )`
  - ▶ `dpx = _mm_or_ps( dpx,dpz );`
  - ▶ `result = _mm_or_ps( dpx,dpw );`
- ▶ Or, with blend:
  - ▶ `tmp = _mm_blend_ps( dpx,dpy, 2 );`
  - ▶ `tmp2 = _mm_blend_ps( dpz,dpw, 8 );`
  - ▶ `result = _mm_blend_ps( tmp, tmp2, 0xc );`

# Assignment

- ▶ Write a program which takes a single command line argument. This will be the name of a two-track (stereo) wave file (some examples are found on the class webpage)
  - ▶ Exclusively using AVX intrinsics, swap the left and right channels
  - ▶ Write the output to a file named “swapped.wav”
  - ▶ For basic credit, support u8, s16, and f32 waves
  - ▶ For bonus [+50%], support s24 waves (again, using AVX intrinsics).
- ▶ Non-SSE program: [chanswap.cpp](#)

# Sources

- ▶ <https://software.intel.com/en-us/articles/intel-software-development-emulator#faq>
- ▶ <http://www.tomshardware.com/reviews/intel-drops-pentium-brand,1832-2.html>
- ▶ <https://www.mathworks.com/matlabcentral/answers/93455-what-is-the-sse2-instruction-set-how-can-i-check-to-see-if-my-processor-supports-it?requestedDomain=www.mathworks.com>
- ▶ <http://softpixel.com/~cwright/programming/simd/ssse3.php>
- ▶ <https://software.intel.com/sites/default/files/m/8/b/8/D9156103.pdf>
- ▶ [https://msdn.microsoft.com/en-us/library/y08s279d\(v=vs.100\).aspx](https://msdn.microsoft.com/en-us/library/y08s279d(v=vs.100).aspx)
- ▶ <http://www.linuxjournal.com/content/introduction-gcc-compiler-intrinsics-vector-processing?page=0,2>
- ▶ <http://en.cppreference.com/w/cpp/language/alignas>
- ▶ <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>
- ▶ MSDN SSE reference
- ▶ <http://ok-cleek.com/blogs/?p=20540>
- ▶ <http://www.liranuna.com/sse-intrinsics-optimizations-in-popular-compilers/>
- ▶ <https://www.cilkplus.org/tutorial-pragma-simd>
- ▶ <https://stackoverflow.com/questions/2804902/whats-the-difference-between-logical-sse-intrinsics>
- ▶ Daniel Kusswurm. Modern X86 Assembly Language Programming. Apress.
- ▶ Peter Kankowski. Implementing strcmp, strlen, and strstr using SSE 4.2 instructions.  
[https://www.strchr.com/strcmp\\_and\\_strlen\\_using\\_sse\\_4.2](https://www.strchr.com/strcmp_and_strlen_using_sse_4.2)

Created using L<sup>A</sup>T<sub>E</sub>X.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>