

# SSE

## Review

- ▶ We've seen SSE intrinsics and datatypes (int, float, double)
- ▶ We've seen float, double, and integer operations
- ▶ Now we'll discuss comparison operations

## Comparison

- ▶ SSE has several comparison functions
  - ▶ Numeric
  - ▶ String
- ▶ First, we look at ones for numeric data

## Compare

- ▶ Compare four floats:
- ▶ `v3 = _mm_cmpeq_ps(v1,v2)`
  - ▶ `v3[0] = 0xffffffff` if `v1[0]==v2[0]`, 0 otherwise
  - ▶ Same for other three slots
- ▶ We also have `cmpge` ( $\geq$ ), `cmpgt` ( $>$ ), `cmple` ( $\leq$ ), `cmplt` ( $<$ ), `cmpne` ( $\neq$ )
- ▶ This gives us an easy way to set a register to all one's: `v = _mm_cmpeq_ps(v,v);`

## Compare

- ▶ We also have integer comparisons:
- ▶ `_mm_cmpeq_epi{8,16,32,64}`
  - ▶ Equality
  - ▶ Sets slots to all-ones or all-zeros
- ▶ `_mm_cmpgt_epi{8,16,32,64}`
  - ▶ Greater-than test
  - ▶ Same idea: Set slots to all-ones or all-zeros
- ▶ `_mm_cmplt_epi{8,16,32,64}`
  - ▶ less-than test
  - ▶ Same idea: Set slots to all-ones or all-zeros

# Strings

- ▶ What about strings? Or arbitrary chunks of data?
- ▶ Two kinds of buffers: Explicit-length vs. implicit length
- ▶ Explicit = We know string length beforehand
  - ▶ Typically, software stores length of string separate from string data
  - ▶ Ex: C++ string type
- ▶ Implicit: We don't know string length without looking for sentinel
  - ▶ C strings: `\0` (null) byte
- ▶ SSE supports both forms

## String Ops

- ▶ Four functions
- ▶ Obtain an index:
  - ▶ `int _mm_cmpestri( __m128i a, int la, __m128i b, int lb, int flags )`
    - ▶ Explicit length
  - ▶ `int _mm_cmpistri( __m128i a, __m128i b, int flags )`
    - ▶ Implicit length: CPU scans for null byte
- ▶ Obtain a mask:
  - ▶ `__m128i _mm_cmpestrm( __m128i a, int la, __m128i b, int lb, int flags )`
  - ▶ `__m128i _mm_cmpistrm( __m128i a, __m128i b, int flags )`
- ▶ We'll tackle the mask ones first

# Flags

- ▶ Flags is a bitwise-or of several values
- ▶ First thing: Need to choose size for comparison
- ▶ Choose one of:
  - ▶ `_SIDD_UBYTE_OPS`: unsigned bytes
  - ▶ `_SIDD_SBYTE_OPS`: signed bytes
  - ▶ `_SIDD_UWORD_OPS`: unsigned shorts
  - ▶ `_SIDD_SWORD_OPS`: signed shorts
- ▶ For character strings, usually use `_SIDD_UBYTE_OPS`
  - ▶ We'll assume that for the following discussion



## Result

- ▶ Choose whether result should be returned as a bit mask or byte mask
  - ▶ If flags is OR'd with `_SIDD_UNIT_MASK`:
    - ▶ If input size is byte: Output 16 bytes, each one 0 or 255
    - ▶ If input size is short: Return 8 shorts, each one 0 or 65535
  - ▶ If flags is OR'd with `_SIDD_BIT_MASK`:
    - ▶ If input size is a byte: Result is 16 bits, with upper bits all zero
    - ▶ If input size is short: Result is 8 bits, with upper bits all zero
- ▶ Following slides refer to output “slot”
  - ▶ Will be either a single bit or a single byte or a single short, depending on flags

## Operation

- ▶ Next, choose a comparison operation and bitwise-OR that into flags
- ▶ Choose one of
  - ▶ `_SIDD_CMP_EQUAL_ANY`,
  - ▶ `_SIDD_CMP_RANGES`,
  - ▶ `_SIDD_CMP_EQUAL_EACH`,
  - ▶ `_SIDD_CMP_EQUAL_ORDERED`

## Flags

- ▶ Easiest one to understand: Flags was OR'd with `_SIDD_CMP_EQUAL_EACH`
- ▶ Then: Slot  $i$  of output is true iff  $a[i] == b[i]$
- ▶ This is a “string equal” operation
- ▶ Note: If  $i$  is past end of both strings, slot  $i$  gets true
  - ▶ If  $i$  is only past end of one string, slot  $i$  gets false

## Flags

- ▶ Suppose instead flags included bitwise OR of `_SIDDCMP_EQUAL_ANY`
- ▶ Slot  $i$  of output is true iff  $b[i]$  is found somewhere in  $a$
- ▶ Bytes past end of either string don't count as equal

## Flags

- ▶ Suppose flags included bitwise OR of `_SIDD_CMP_EQUAL_ORDERED`
  - ▶ Slot 0 of output is true iff a appears at start of b
  - ▶ Slot 1 of output is true iff a appears at b[1:]
  - ▶ Slot 2 of output is true iff a appears at b[2:]
  - ▶ ...
- ▶ For all slots past end, set slot to false

## Flags

- ▶ Suppose flags included bitwise OR of `_SIDD_CMP_RANGES`
- ▶ Consider `a` to be a list of character pairs
- ▶ Return true for each slot `i` where `b[i]` is within any of the ranges defined by the character pairs in `a`
- ▶ This will make more sense when we see an example...

## Other Flags

- ▶ There are some other tweaks that can be made
  - ▶ If flags includes `_STDD_NEGATIVE_POLARITY`: Invert results
  - ▶ If flags includes `_SIDD_MASKED_NEGATIVE_POLARITY`: Invert results for indices `0,1,2,...,len(b)`

# Wow

- ▶ Wow! That's a lot to absorb
- ▶ How can we use instructions like this?
- ▶ Let's see some examples



# Example

- ▶ String equality
- ▶ Example program:

```
alignas(16) char aa[] = "pack my box with";
alignas(16) char bb[] = "get thy box\0\0\0\0";
int la = strlen(aa);
int lb = strlen(bb);
__m128i a = _mm_load_si128((__m128i*)aa);
__m128i b = _mm_load_si128((__m128i*)bb);
__m128i res = _mm_cmpestrm(a,la,b,lb,
    _SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_UNIT_MASK );
alignas(16) char rr[16];
_mm_store_si128((__m128i*)rr,res);
```

## Result

- ▶ Using flags = `_SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_UNIT_MASK`

- ▶ Result:

```
a=pack my box with  
b=get thy box  
  FFFFFFFTTTTTFFFFF
```

- ▶ Notice: true wherever strings are equal
  - ▶ The 'T' represents "this slot of output contains eight one bits"
  - ▶ A 'F' means "this slot of output contains eight zero bits"

## Result

- ▶ Using flags = \_SIDD\_UBYTE\_OPS | \_SIDD\_CMP\_EQUAL\_EACH | \_SIDD\_UNIT\_MASK
- ▶ Result:  
a=pack my bag wit  
b=get zat box  
FFFFFFFFTTFFFFFFFF
- ▶ Notice that when we're past end of string, compares as true

## Result

- ▶ Using flags = `_SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_BIT_MASK`

- ▶ Result:

a=pack my box with

b=get thy box

c0,7,0,0,0,0,0,0,0,0,0,0,0,0,0,0

- ▶ Here, the output is not just 0 or 255 in each slot
  - ▶ Take low 16 bits of result: 0xc0, 0x7, 0, 0 (little endian)
  - ▶ That's 0x7c0 = 0b0000 0111 1100 0000
  - ▶ Convert to binary (msb corresponds to end of strings):

000000111110000000

htiw xob ym kcap

xob yht teg

## Result

- ▶ Using flags = `_SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_ANY | _SIDD_UNIT_MASK`
- ▶ Result:  
a=pack my bag wit  
b=get zat box  
101101111000000
- ▶ Since a contains 'g', 't', ' ', 'b', those output slots are true
- ▶ Since a does not contain 'e', 'z', 'o', or 'x', those slots are false
- ▶ Slots corresponding to locations past end of either string are false

## Result

- ▶ Using flags = \_SIDD\_UBYTE\_OPS | \_SIDD\_CMP\_EQUAL\_ORDERED | \_SIDD\_UNIT\_MASK
- ▶ Result:  
a=my bag  
b=pack my bag, sir  
0000010000000000
- ▶ Notice the '1' appears where 'a' first appears in b

## Result

- ▶ Using flags = \_SIDD\_UBYTE\_OPS | \_SIDD\_CMP\_EQUAL\_ORDERED | \_SIDD\_UNIT\_MASK
- ▶ Result:  
a=my bag  
b=my bag! my bag!!  
1000000010000000
- ▶ Since two matches, two 1's

## Result

- ▶ Using flags = `_SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_ORDERED | _SIDD_UNIT_MASK`
- ▶ Result:  
a=foofoo  
b=foofoofoofoo  
1001001000000000
- ▶ Overlapping matches are detected as separate matches



## Result

- ▶ Using flags = `_SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_ORDERED | _SIDD_UNIT_MASK`
- ▶ Result:  
a=my bag!  
b=my bag is mine  
000000000000000000
- ▶ Since the ! does not appear in b, no match

## Result

- ▶ Using `flags = _SIDD_UBYTE_OPS | _SIDD_CMP_RANGES | _SIDD_UNIT_MASK`
- ▶ Result:  

```
a=am  
b=pack my box with  
  0111010010000101
```
- ▶ Returns true for each slot where character from b is in {a,b,c,...,l,m}

## Result

- ▶ Using flags = \_SIDD\_UBYTE\_OPS | \_SIDD\_CMP\_RANGES |  
\_SIDD\_UNIT\_MASK

- ▶ Result:

a=aemmwz

b=pack my box with

0110011010101000

- ▶ Returns true where any character of b is in {a,b,c,d,e} or {m} or  
{w,x,y,z}

# Index

- ▶ What about the index functions?
  - ▶ `int _mm_cmpestri( __m128i a, int la, __m128i b, int lb, int flags )`
  - ▶ `int _mm_cmpistri( __m128i a, __m128i b, int flags )`
- ▶ Operation is similar to mask, but returns an index
  - ▶ If flags includes `_SIDD_LEAST_SIGNIFICANT`: Return index of first place where comparison succeeded
  - ▶ If flags includes `_SIDD_MOST_SIGNIFICANT`: Return index of last place where comparison succeeded

## Example

- ▶ Using `int res = _mm_cmpestri( ... , _SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_LEAST_SIGNIFICANT);`
- ▶ Result:  
a=crack a boxcar  
b=pack my box with  
res=7
- ▶ Match starts at character 7 (the space character before 'box')

## Example

- ▶ Using `int res = _mm_cmpestri(..., _SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_MOST_SIGNIFICANT);`
- ▶ Result:  
a=crack a boxcar  
b=pack my box with  
res=10
- ▶ Last match is at index 10

## Example

- ▶ Using `int res = _mm_cmpestri( ... , _SIDD_UBYTE_OPS | _SIDD_CMP_EQUAL_EACH | _SIDD_MOST_SIGNIFICANT);`
- ▶ Result:  
a=crackerjack  
b=pack my box with  
res=16
- ▶ 16 because no match (we're using EQUAL\_EACH here)

## Uses

- ▶ EQUAL\_EACH is like strcmp()
- ▶ EQUAL\_ANY is like strspn()
- ▶ EQUAL\_ORDERED is like strstr()
- ▶ EQUAL\_RANGES doesn't have a direct C equivalent



# Assignment

- ▶ Often, sound recordings contain noise in areas that should be silent
- ▶ Write a program to fix this using SSE:
  - ▶ It should take two command line arguments: The filename of a WAV file and a silence threshold
  - ▶ Read the WAV file. Any values within “threshold” of silence should be squelched to silence
    - ▶ For u8: samples are in range 0...255, so silence == 127
    - ▶ For s16, samples are in range -32768...32767, so silence == 0
  - ▶ Write output to the file “silenced.wav”
- ▶ Your program should work with u8 and s16 WAV's
- ▶ Example code: [silence.cpp](#)

# Sources

- ▶ <https://software.intel.com/en-us/articles/intel-software-development-emulator#faq>
- ▶ <http://www.tomshardware.com/reviews/intel-drops-pentium-brand,1832-2.html>
- ▶ <https://www.mathworks.com/matlabcentral/answers/93455-what-is-the-sse2-instruction-set-how-can-i-check-to-see-if-my-processor-supports-it?requestedDomain=www.mathworks.com>
- ▶ <http://softpixel.com/cwright/programming/simd/ssse3.php>
- ▶ <https://software.intel.com/sites/default/files/m/8/b/8/D9156103.pdf>
- ▶ [https://msdn.microsoft.com/en-us/library/y08s279d\(v=vs.100\).aspx](https://msdn.microsoft.com/en-us/library/y08s279d(v=vs.100).aspx)
- ▶ <http://www.linuxjournal.com/content/introduction-gcc-compiler-intrinsics-vector-processing?page=0,2>
- ▶ <http://en.cppreference.com/w/cpp/language/alignas>
- ▶ <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>
- ▶ MSDN SSE reference
- ▶ <https://wiki.multimedia.cx/index.php/YUV4MPEG2>
- ▶ <http://ok-cleek.com/blogs/?p=20540>
- ▶ <http://www.liranuna.com/sse-intrinsics-optimizations-in-popular-compilers/>
- ▶ <https://www.cilkplus.org/tutorial-pragma-simd>
- ▶ A. Ortiz. "Teaching the SIMD Execution Model: Assembling a Few Parallel Programming Skills." Proceedings of 2003 ACM SIGCSE.
- ▶ Nasm documentation.
- ▶ Greggo. x86 - SSE Compare Packed Unsigned Bytes.  
<https://stackoverflow.com/questions/16204663/sse-compare-packed-unsigned-bytes>
- ▶ John D. Cook. Converting color to grayscale. <https://www.johndcook.com/blog/2009/08/24/algorithms-convert-color-grayscale/>

Created using L<sup>A</sup>T<sub>E</sub>X.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>