

# SSE

## Motivation

- ▶ Sometimes, SIMD doesn't quite provide what we want
- ▶ We need to synthesize the desired operation out of the building blocks we *do* have.

# NOT

- ▶ Suppose we have image data, as RGBA
- ▶ Invert all colors of an image: There's no bitwise NOT operation
- ▶ How can we accomplish this?

# XOR

- ▶ We can press XOR into service

```
__m256* p = (__m256i*) img.pixels();  
__m256i ones = _mm256_cmpeq_epi8(ones,ones);  
for(unsigned i=0;i<n;i+=32,p++){  
    v = _mm256_lddqu_si256(p);  
    v = _mm256_xor_si256(v,ones);  
    _mm256_storeu_si256(p,v);  
}
```

- ▶ But there's a problem. What is it?

# Problem

- ▶ This also inverts the *alpha* channel
  - ▶ That converts image to be entirely transparent.
  - ▶ We want to just invert the RGB channels.
- ▶ How to do this?

## Solution

- Change our mask:

```
__m256* p = (__m256i*) img.pixels();
__m256i ones = _mm256_set_epi8( 255,255,255,0, 255,255,255,0,
    255,255,255,0, ... );
for(unsigned i=0;i<n;i+=32,p++){
    v = _mm256_lddqu_si256(p);
    v = _mm256_xor_si256(v,ones);
    _mm256_storeu_si256(p,v);
}
```

# Averaging

- ▶ Suppose we want to convert image to greyscale
- ▶ Several ways to accomplish this
- ▶ Option 1: Replicate one channel to all slots
  - ▶ Human eye most sensitive to green, so that's a logical choice
- ▶ This is just a swizzle operation

## Code

```
uint8_t tmp[32] = {  
    1,1,1,3,    5,5,5,7,    9,9,9,11,    13,13,13,15,  
    1,1,1,3,    5,5,5,7,    9,9,9,11,    13,13,13,15,  
};  
__m256i shuf = _mm256_lddqu_si256( (__m256i*) tmp);  
__m256i* p = (__m256i*) img.pixels();  
for(unsigned i=0;i<n;i+=32,p++){  
    __m256i v = _mm256_load_si256(p);  
    v = _mm256_shuffle_epi8(v, shuf );  
    _mm256_storeu_si256(p,v);  
}
```

- ▶ Timings: SIMD: 6009  $\mu$ sec; non-SIMD: 20967  $\mu$ sec



## Problem

- ▶ What if we have  $r=255$ ,  $g=0$ ,  $b=255$  (bright magenta)?
  - ▶ This scheme will output black

## Strategy 2

- ▶ What if we want to take average?
- ▶ Compute:  $(r+g+b)/3$
- ▶ Or: Could compute:  $r/3 + g/3 + b/3$
- ▶ Which approach should we use?

## Ex

- ▶ Suppose  $r=2, g=2, b=2$ .
- ▶ Then  $r/3+g/3+b/3 = 0+0+0 = 0$
- ▶ What about  $(r+g+b)/3$ ?
  - ▶  $(2+2+2)/3 = 6/3 = 2$
- ▶ So it looks like we want  $(r+g+b)/3$
- ▶ But...We have a problem. What is it?

## Problem

- ▶ Suppose  $r=255$ ,  $g=255$ ,  $b=255$
- ▶ If we add with saturation: We get  $(255+255+255) = 255$ 
  - ▶ Then  $255/3 = 85$
- ▶ If add without saturation:  $255+255+255 = 253$ 
  - ▶  $255+255=510 \rightarrow 254$
  - ▶ Then  $254+255=509 \rightarrow 253$
  - ▶ And we get  $253/3 = 84$
- ▶ How to solve this problem?

## Solution?

- ▶ Promote bytes to int, then divide by 3
- ▶ Unfortunately, there's no SSE integer divide instruction
  - ▶ Only floating point divide
- ▶ Darn.

## Method

- ▶ Maybe we can use another method: *Lightness*
- ▶ Take max of r,g,b and min of r,g,b
- ▶ Sum those
- ▶ Divide by two
- ▶ But...Didn't we just decide there was no division operator?

## Division

- ▶ Yes, but there's an instruction that applies concept of *strength reduction*
  - ▶ Division by 2 == Shift right one place
- ▶ `_mm256_avg_epu8(a,b)`: Compute  $(a+b) \gg 1$

## Step 1

- ▶ Load eight pixels to a register

```
__m256i v1 = _mm256_lddqu_si256(p);
```

$$v1 = \begin{array}{c} (hi) \alpha_7, b_7, g_7, r_7, \alpha_6, b_6, g_6, r_6, \alpha_5, b_5, g_5, r_5, \alpha_4, b_4, g_4, r_4 \\ \alpha_3, b_3, g_3, r_3, \alpha_2, b_2, g_2, r_2, \alpha_1, b_1, g_1, r_1, \alpha_0, b_0, g_0, r_0 (low) \end{array} |$$

- ▶ First byte in memory (red, pixel 0) goes to slot 0 of YMM; last byte in memory (alpha, pixel 7) goes to slot 31 of YMM
- ▶ The | represents the lane division



## Step 2

- Do a 32-bit right shifts on v1:

```
__m256i v2 = _mm256_srli_epi32( v1, 8 );
```

```
__m256i v3 = _mm256_srli_epi32( v1, 16 );
```

$$v2 = \begin{array}{l} 0, \alpha_7, b_7, g_7, 0, \alpha_6, b_6, g_6, 0, \alpha_5, b_5, g_5, 0, \alpha_4, b_4, g_4 | \\ 0, \alpha_3, b_3, g_3, 0, \alpha_2, b_2, g_2, 0, \alpha_1, b_1, g_1, 0, \alpha_0, b_0, g_0 \end{array}$$

$$v3 = \begin{array}{l} 0, 0, \alpha_7, b_7, 0, 0, \alpha_6, b_6, 0, 0, \alpha_5, b_5, 0, 0, \alpha_4, b_4 | \\ 0, 0, \alpha_3, b_3, 0, 0, \alpha_2, b_2, 0, 0, \alpha_1, b_1, 0, 0, \alpha_0, b_0 \end{array}$$

## Step 3

- ▶ Take the max's and mins

```
max = _mm256_max_epu8(v1,v2);
```

```
max = _mm256_max_epu8(max,v3);
```

```
min = _mm256_min_epu8(v1,v2);
```

```
min = _mm256_min_epu8(min,v3);
```

- ▶ slots 0, 4, 8, 12, ... , 28 of “max” have maximums of (r,g,b) for the eight pixels
- ▶ slots 0, 4, 8, 12, ... , 28 of “min” have minimums of (r,g,b) for the eight pixels
- ▶ slots 3, 7, 11, ..., 31 of “max” have original  $\alpha$  values (since we always max'd with 0)

## Step 4

- ▶ Take the average:

```
avg = _mm256_avg_epu8( max,min );
```

## Step 5

- ▶ Mask the parts where  $\alpha$  will go:
- ▶ Let mask = {255 (byte 0),255,255,0,255,255,255,0,...,255,255,255,0 (byte 31)}
- ▶ `avg = _mm256_and_si256( avg, mask )`
- ▶ Mask the  $\alpha$  from the original
- ▶ `alpha = _mm256_andnot_si256( mask, v1 );`
- ▶ Combine:
- ▶ `avg = _mm256_or_si256( avg, alpha );`

## Shuffle

- ▶ Now we can distribute the values back to their respective pixel locations
- ▶ Let `shuf = {0,0,0,3, 4,4,4,7, 8,8,8,11, 12,12,12,15, 0,0,0,3, 4,4,4,7, 8,8,8,11, 12,12,12,15};`
  - ▶ Byte 0 of `shuf`=0; byte 31 of `shuf`=15
- ▶ Then:  
`res = _mm256_shuffle_epi8( avg, shuf );`

# Timing

- ▶ For 3648x2736 image, -O3:
  - ▶ SIMD: 5617  $\mu$ sec
  - ▶ Non-SIMD: 21612  $\mu$ sec

## Option

- ▶ Yet another option: Take account of qualities of human perception
  - ▶ Eye is more sensitive to green than red and more sensitive to red than blue
- ▶ NTSC method:  $0.21r + 0.72g + 0.07b$

## Option

- ▶ To avoid floating point conversions, we can alter this a bit:
  - ▶  $\frac{1}{4} \cdot r \rightarrow 0.25$
  - ▶  $\frac{11}{16} \cdot g \rightarrow 0.6875$
  - ▶  $\frac{1}{16} \cdot b \rightarrow 0.0625$
- ▶ Divide by 4: Shift right 2
- ▶ Divide by 16: Shift right by 4
- ▶ But what about 11/16?



## Fraction

- ▶ Could divide by 16 then multiply by 11
  - ▶ But: This loses precision
- ▶ Ex: Suppose green = 111
  - ▶ Correct result:  $11/16 * 100 = 76.3125 = 76$
  - ▶ If we divide by 16, we get 6; multiply by 11 gives 66
  - ▶ Error: 13%
- ▶ Could multiply by 11 and then divide by 16
  - ▶ But: This overflows whenever green > 23
- ▶ Ideas...?

# Idea

- ▶ One way:
  - ▶ Observe:  $\frac{11}{16} = \frac{8+2+1}{16} = \frac{8}{16} + \frac{2}{16} + \frac{1}{16} = \frac{1}{2} + \frac{1}{8} + \frac{1}{16}$
  - ▶ So we could compute  $g/2 + g/8 + g/16$
  - ▶ Thus:
    - ▶  $\text{tmp1} = g \gg 1$
    - ▶  $\text{tmp2} = g \gg 3$
    - ▶  $\text{tmp3} = g \gg 4$
    - ▶  $\text{Result} = \text{tmp1} + \text{tmp2} + \text{tmp3}$
  - ▶ Ex: For 111: We get  $55+13+8 = 74$  (error = 2.6%)

## Note

- ▶ If we wanted to be more exact, we could use
  - ▶  $\frac{7}{32} \cdot r = 0.21875 \cdot r$  (compared to 0.21\*red for NTSC)
  - ▶  $\frac{23}{32} \cdot g = 0.71875 \cdot g$  (compared to 0.72\*green)
  - ▶  $\frac{2}{32} \cdot b = 0.0625 \cdot b$  (compared to 0.07\*blue)

## Note

- ▶ This would be a bit more work:
  - ▶  $\left(\frac{4}{32} + \frac{2}{32} + \frac{1}{32}\right) r = \left(\frac{1}{8} + \frac{1}{16} + \frac{1}{32}\right) r = (r \gg 3) + (r \gg 4) + (r \gg 5)$
  - ▶  $\left(\frac{16}{32} + \frac{4}{32} + \frac{2}{32} + \frac{1}{32}\right) g = \left(\frac{1}{2} + \frac{1}{8} + \frac{1}{16} + \frac{1}{32}\right) g = (g \gg 1) + (g \gg 3) + (g \gg 4) + (g \gg 5)$
  - ▶  $\left(\frac{2}{32}\right) b = \left(\frac{1}{16}\right) b = b \gg 4$
- ▶ We can choose whether we want to trade off speed for accuracy

## Problem

- ▶ There's no 8-bit integer shift operation in AVX (or SSE, for that matter)
- ▶ There are 16, 32, and 64 bit integer shifts
- ▶ But we can build an 8-bit shift out of a larger shift + mask operation

## Consider

- ▶ Suppose we have a YMM register with a series of bytes  $b_0, b_1, \dots$  and we want to shift each byte to the right by 3

- ▶ Register contents:  
$$\begin{array}{cccccccc} b_{31,7} & b_{31,6} & b_{31,5} & b_{31,4} & b_{31,3} & b_{31,2} & b_{31,1} & b_{31,0} \\ b_{30,7} & b_{30,6} & b_{30,5} & b_{30,4} & b_{30,3} & b_{30,2} & b_{30,1} & b_{30,0} \\ & & & & \vdots & & & \\ b_{0,7} & b_{0,6} & b_{0,5} & b_{0,4} & b_{0,3} & b_{0,2} & b_{0,1} & b_{0,0} \end{array}$$

- ▶ Notation:  $b_{i,j}$  = bit  $j$  of byte  $i$

## Shift

- ▶ Suppose we do a 16-bit integer logical right shift:  
\_mm256\_srli\_epi16( v, 3)
- ▶ Register contents before shift:

$$\begin{array}{cccccccc} b_{31,7} & b_{31,6} & b_{31,5} & b_{31,4} & b_{31,3} & b_{31,2} & b_{31,1} & b_{31,0} \\ b_{30,7} & b_{30,6} & b_{30,5} & b_{30,4} & b_{30,3} & b_{30,2} & b_{30,1} & b_{30,0} \\ & & & \vdots & & & & \\ b_{1,7} & b_{1,6} & b_{1,5} & b_{1,4} & b_{1,3} & b_{1,2} & b_{1,1} & b_{1,0} \\ b_{0,7} & b_{0,6} & b_{0,5} & b_{0,4} & b_{0,3} & b_{0,2} & b_{0,1} & b_{0,0} \end{array}$$

## Shift

$$\begin{array}{c} 000b_{31,7}b_{31,6}b_{31,5}b_{31,4}b_{31,3} \\ b_{31,2}b_{31,1}b_{31,0}b_{30,7}b_{30,6}b_{30,5}b_{30,4}b_{30,3} \\ \vdots \end{array}$$

- ▶ After shift:

$$\begin{array}{c} 000b_{1,7}b_{1,6}b_{1,5}b_{1,4}b_{1,3} \\ b_{1,2}b_{1,1}b_{1,0}b_{0,7}b_{0,6}b_{0,5}b_{0,4}b_{0,3} \end{array}$$

- ▶ Bytes 1,3,5,... are OK: They got zeros shifted in, like we wanted
- ▶ But byte 0 contains some “spillover” from byte 1
  - ▶ Likewise for bytes 2, 4, 6, ...



## Mask

- ▶ Suppose we have a mask:  $0b0001\ 1111 = 0x1f$  in all 32 positions
- ▶ We could then AND our result with that mask
- ▶ This would give us the result we want for an unsigned byte-shift by 3

## Code

- ▶ We can write a function to help us out:

```
__m256i shiftRightByte(__m256i v, unsigned count){  
    v = _mm256_srli_epi16( v, count);  
    __m256i mask = _mm256_set1_epi8( 0xff >> count );  
    return _mm256_and_si256( v, mask );  
}
```

## What We Need

- ▶  $\frac{1}{4} \cdot r \rightarrow$  Shift right 2
- ▶  $\frac{11}{16} \cdot g \rightarrow \frac{1}{2}g + \frac{1}{8}g + \frac{1}{16}g \rightarrow$  Shift 1, shift 3, shift 4
- ▶  $\frac{1}{16} \cdot b \rightarrow$  Shift right 4

## Problem

- ▶ We seem to have another problem: We have different operations to apply to red, green, and blue channels
- ▶ SIMD is happiest when we can apply uniform operations across the whole register
- ▶ We'll deal with this next time!

# Assignment

- ▶ None!
  - ▶ Just get caught up on the existing labs...

# Sources

- ▶ <https://software.intel.com/en-us/articles/intel-software-development-emulator#faq>
- ▶ <http://www.tomshardware.com/reviews/intel-drops-pentium-brand,1832-2.html>
- ▶ <https://www.mathworks.com/matlabcentral/answers/93455-what-is-the-sse2-instruction-set-how-can-i-check-to-see-if-my-processor-supports-it?requestedDomain=www.mathworks.com>
- ▶ <http://softpixel.com/cwright/programming/simd/ssse3.php>
- ▶ <https://software.intel.com/sites/default/files/m/8/b/8/D9156103.pdf>
- ▶ [https://msdn.microsoft.com/en-us/library/y08s279d\(v=vs.100\).aspx](https://msdn.microsoft.com/en-us/library/y08s279d(v=vs.100).aspx)
- ▶ <http://www.linuxjournal.com/content/introduction-gcc-compiler-intrinsics-vector-processing?page=0,2>
- ▶ <http://en.cppreference.com/w/cpp/language/alignas>
- ▶ <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>
- ▶ MSDN SSE reference
- ▶ <https://wiki.multimedia.cx/index.php/YUV4MPEG2>
- ▶ <http://ok-cleek.com/blogs/?p=20540>
- ▶ <http://www.liranuna.com/sse-intrinsics-optimizations-in-popular-compilers/>
- ▶ <https://www.cilkplus.org/tutorial-pragma-simd>
- ▶ A. Ortiz. "Teaching the SIMD Execution Model: Assembling a Few Parallel Programming Skills." Proceedings of 2003 ACM SIGCSE.
- ▶ Nasm documentation.
- ▶ Greggo. x86 - SSE Compare Packed Unsigned Bytes.  
<https://stackoverflow.com/questions/16204663/sse-compare-packed-unsigned-bytes>
- ▶ John D. Cook. Converting color to grayscale. <https://www.johndcook.com/blog/2009/08/24/algorithms-convert-color-grayscale/>

Created using L<sup>A</sup>T<sub>E</sub>X.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>